



NoSQL

Alexey Zinovyev, Java/BigData Trainer in EPAM





With IT since 2007
With Java since 2009
With Hadoop since 2012
With EPAM since 2015

About

Contacts

E-mail : Alexey_Zinovyev@epam.com

Twitter : @zaleslaw @BigDataRussia

Facebook: <https://www.facebook.com/zaleslaw>

vk.com/big_data_russia Big Data Russia

vk.com/java_jvm Java & JVM langs

The Good Old Days



PT700 53.8

One of these fine days...



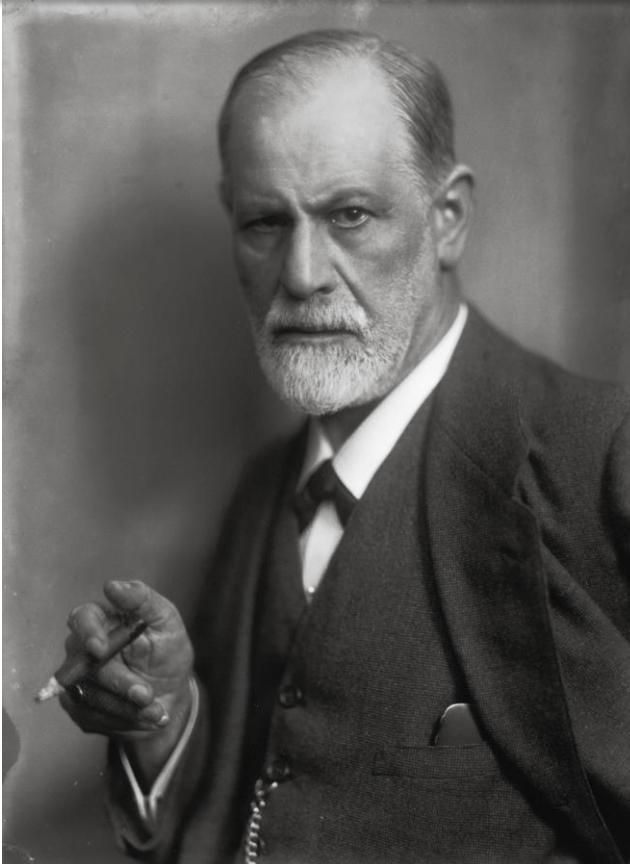
We have a NoSQL job for you, son!



But you like SQL and HATE nontraditional data



Let's talk about it, boy...



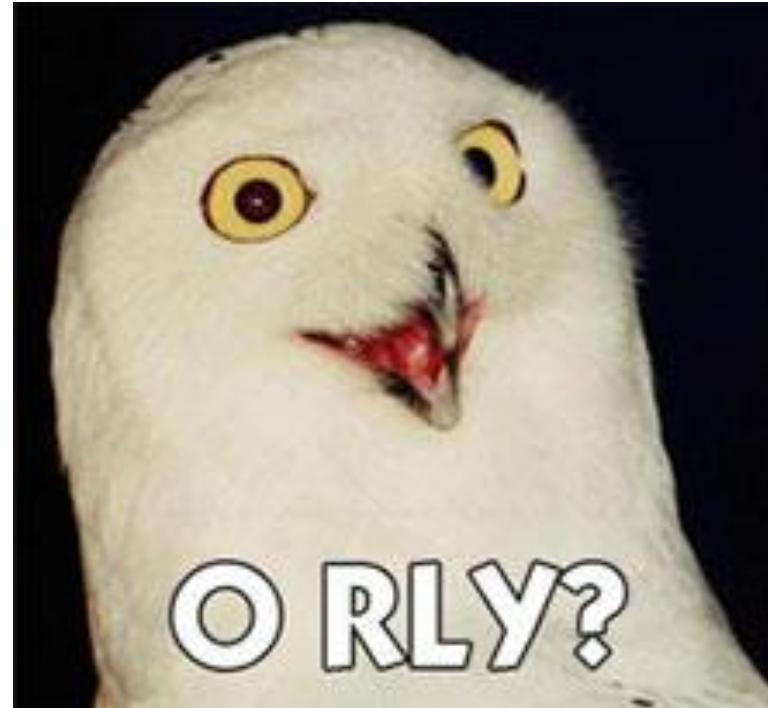
The Database Market



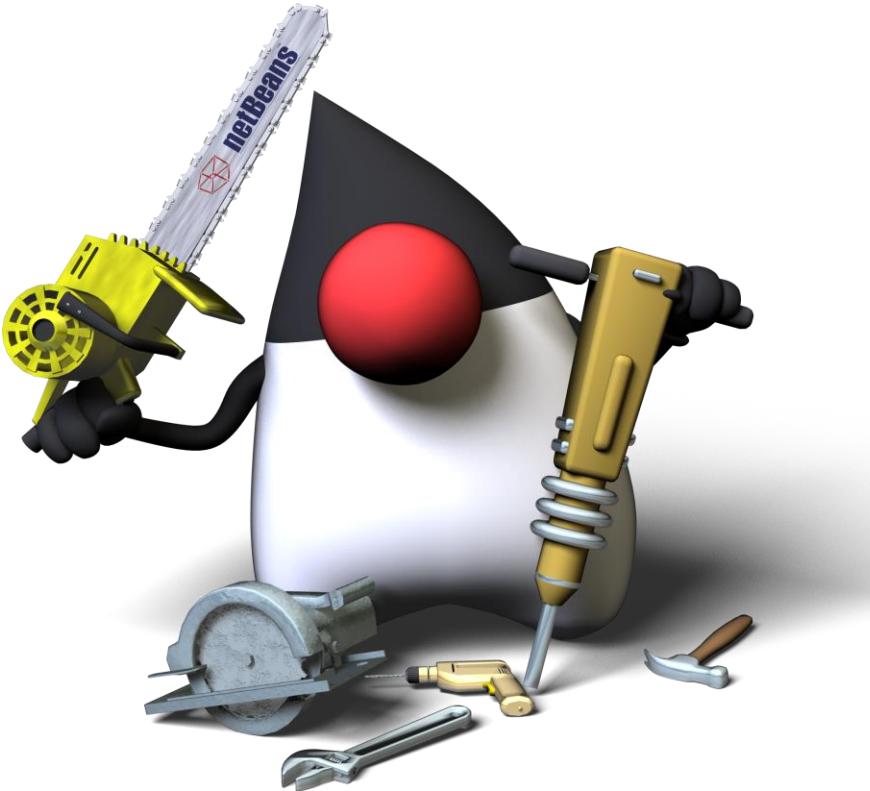
Let's use Cassandra to keep logs ONLY ..



**For logs
ONLY?**



Case #0 : PosgreSQL & Cassandra



A black and white aerial photograph of the New York City skyline. In the foreground, the Hudson River is visible with several industrial buildings and smokestacks emitting plumes of smoke. The midground shows a dense cluster of skyscrapers, including the Art Deco style Chrysler Building and the Art Deco style Empire State Building. The background features a vast, cloudy sky.

Backend Development in 1997

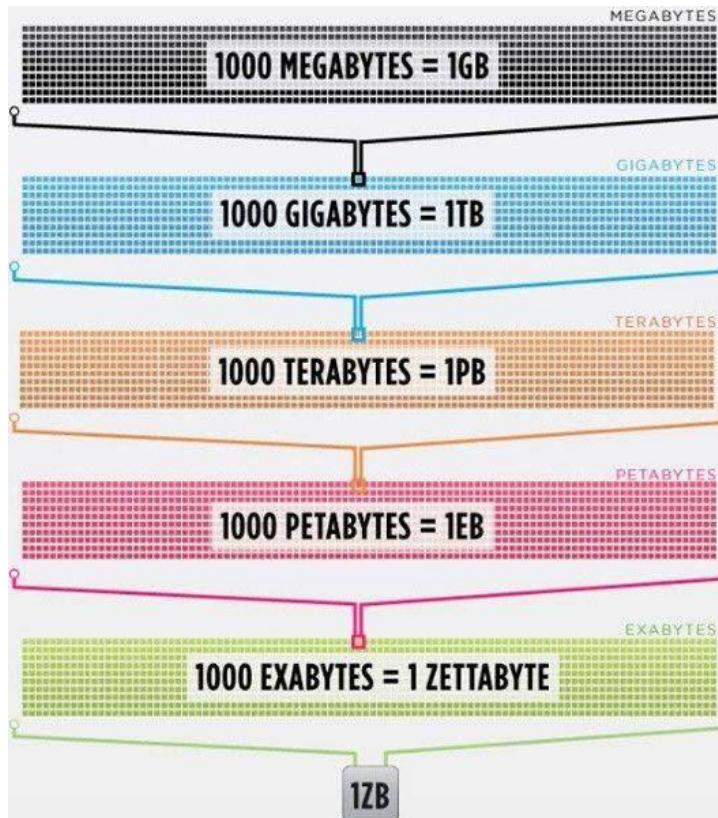


Backend Development in 2017

**10⁶
rows in
MySQL**



GB->TB->PB->?



This chapter covers ...

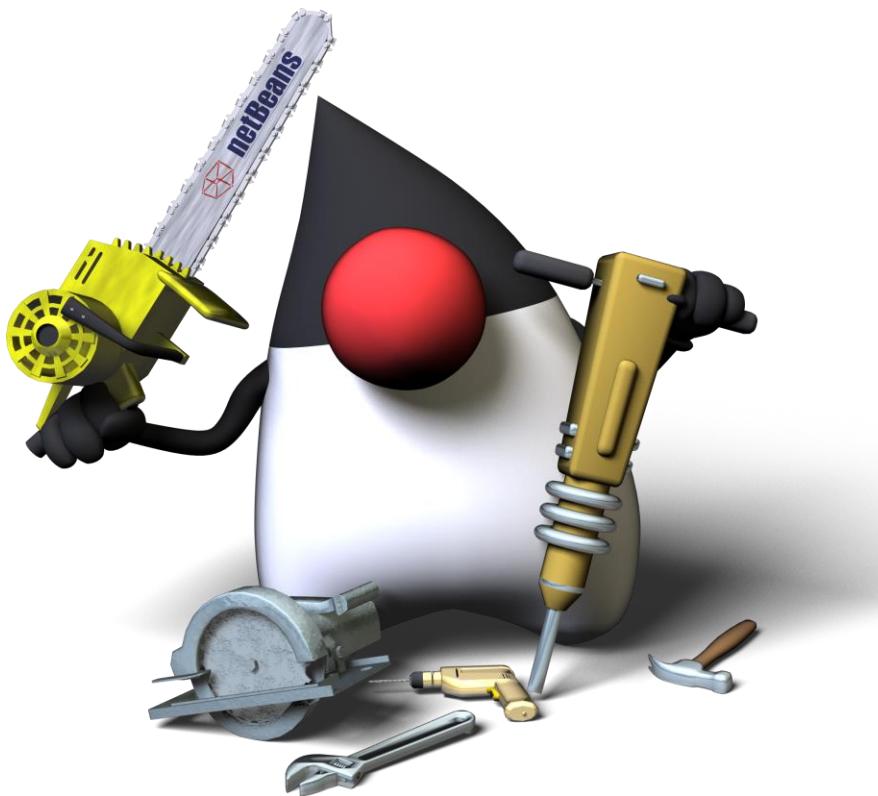
- Highload applications
- Sharding & replication
- NoSQL databases

Case #1 : Likes in Classmates

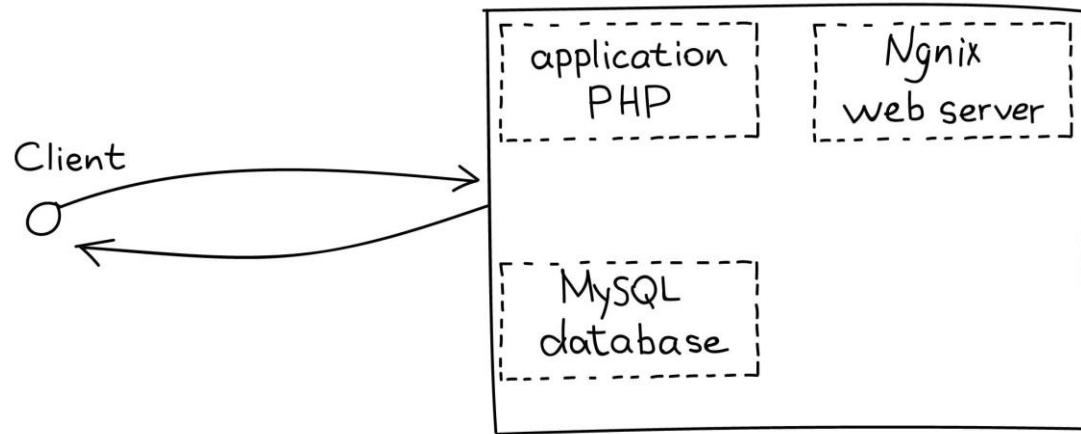


HIGHLOAD APPLICATIONS

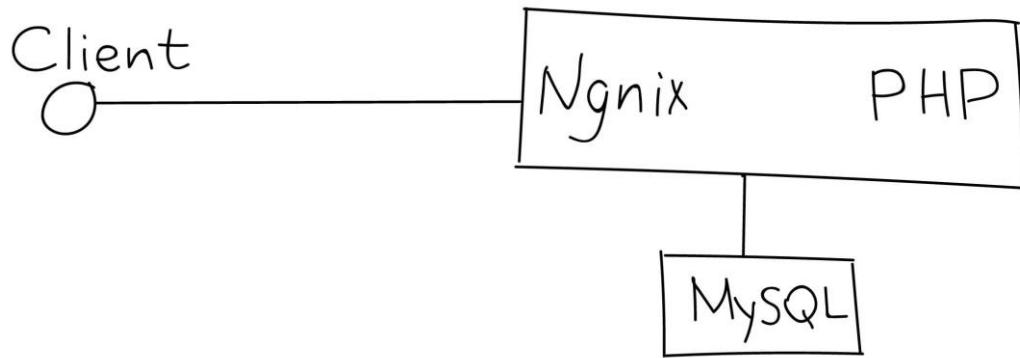
Case #2 : Pet E-Commerce



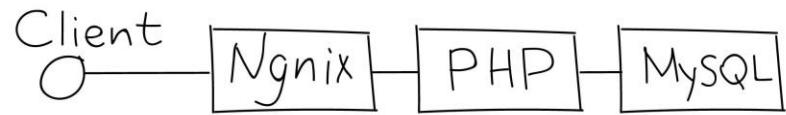
Simple web application #1



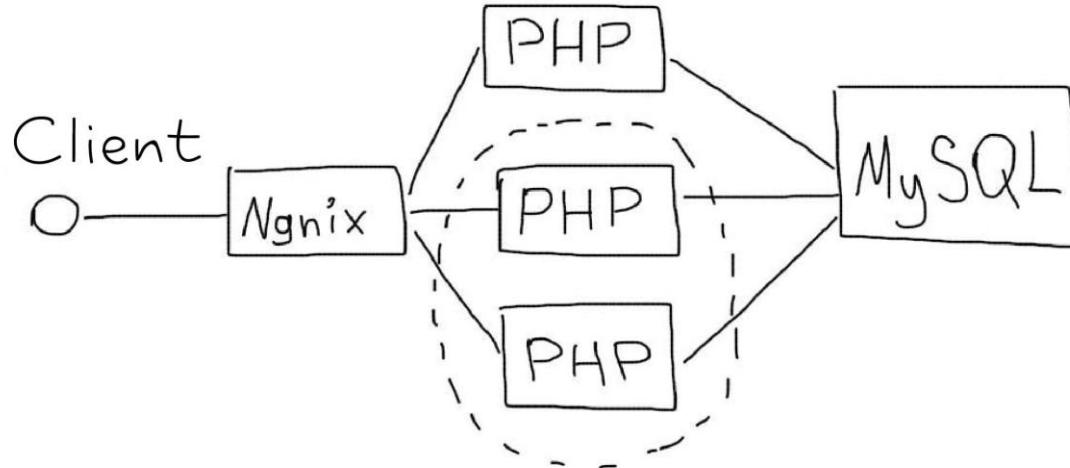
Simple web application #2



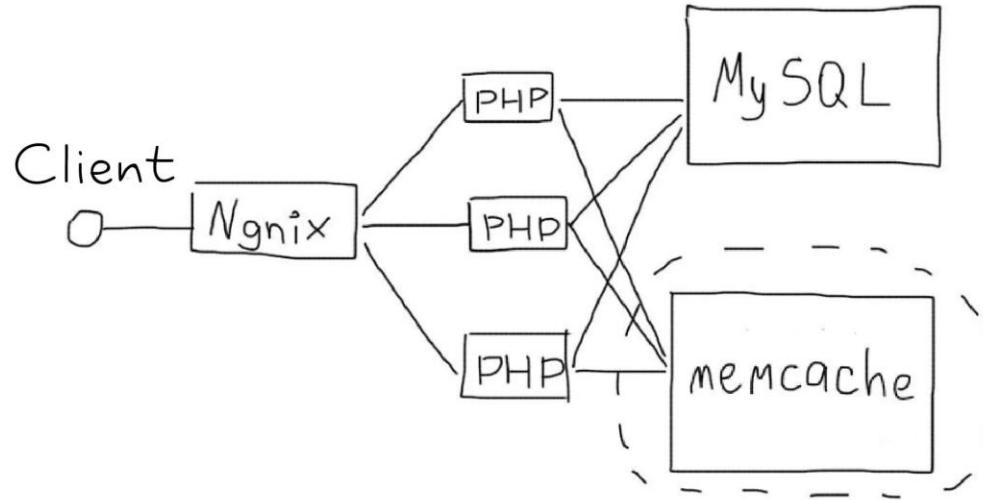
Simple web application #3



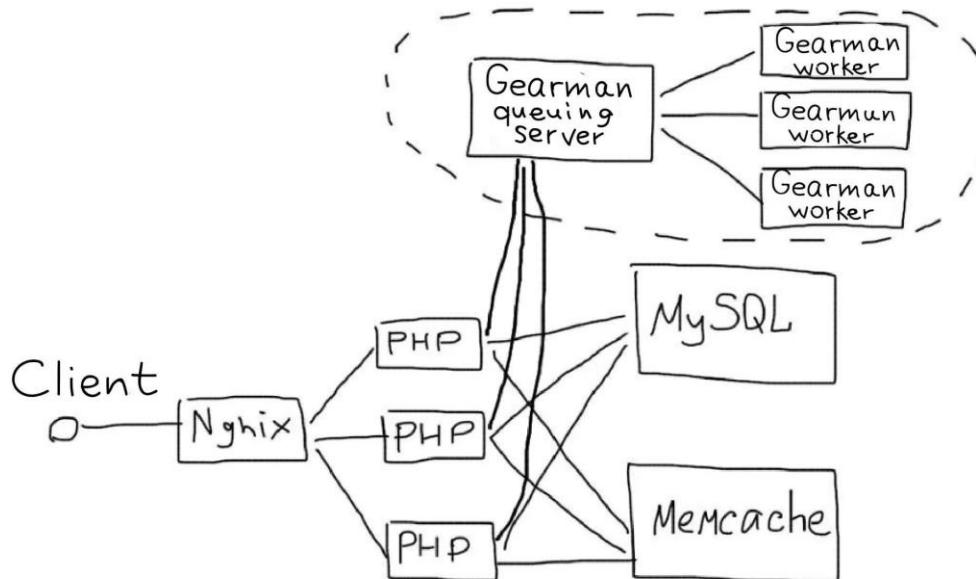
Simple web application #4



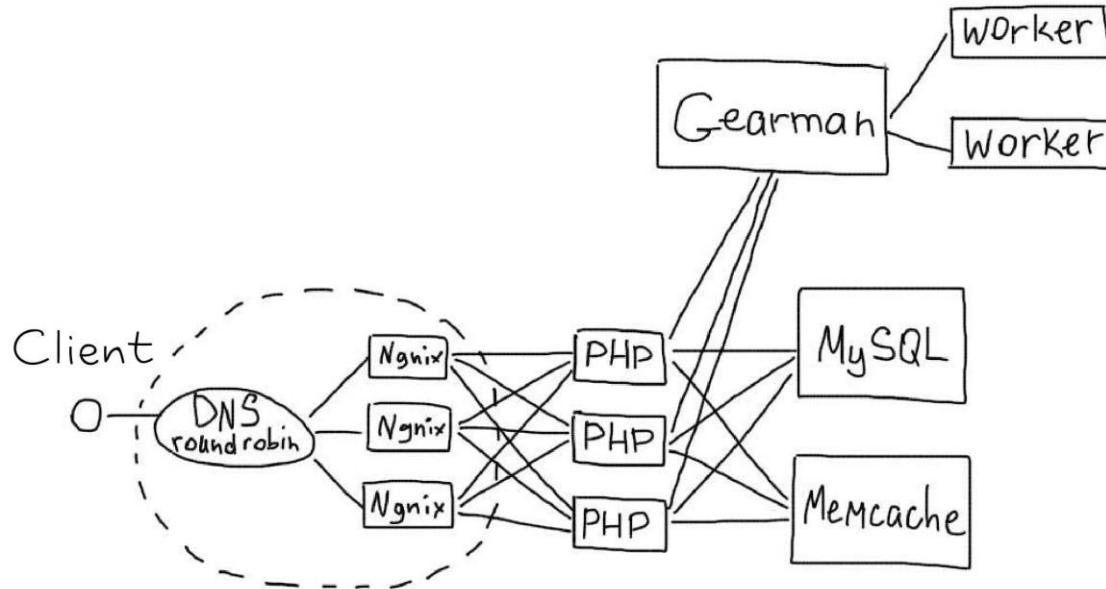
Simple web application #5



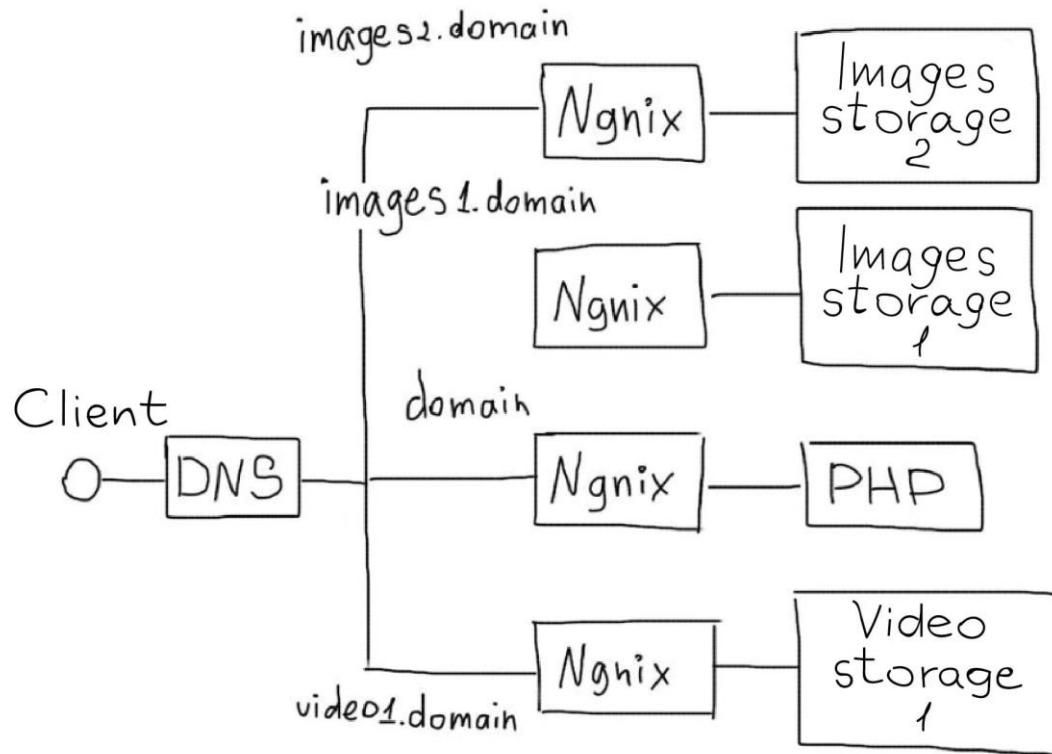
Simple web application #6



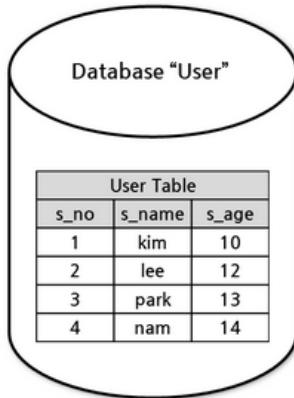
Simple web application #7



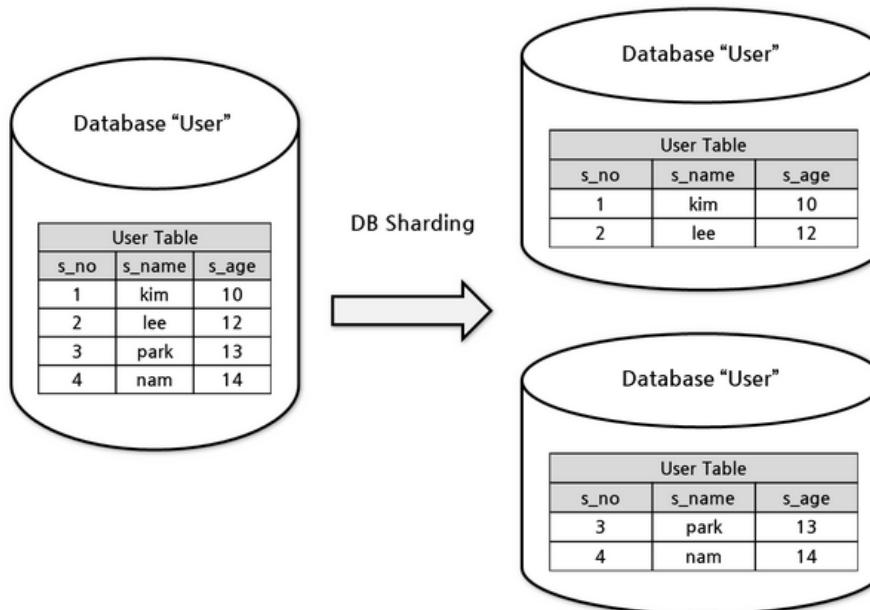
Simple web application #8



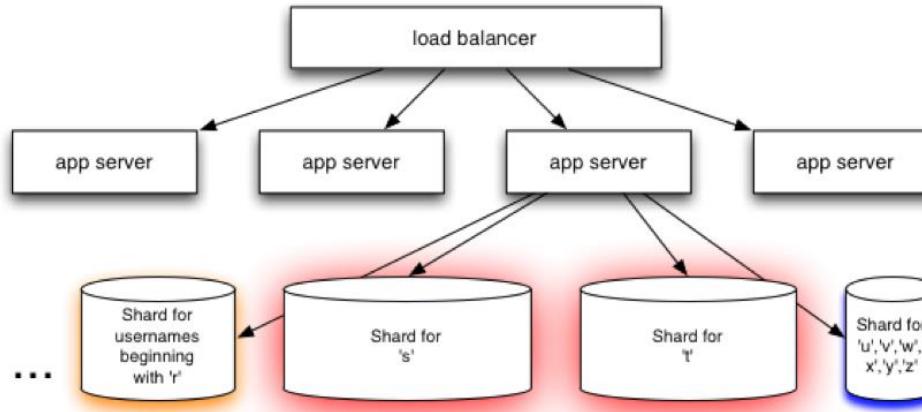
Big Table



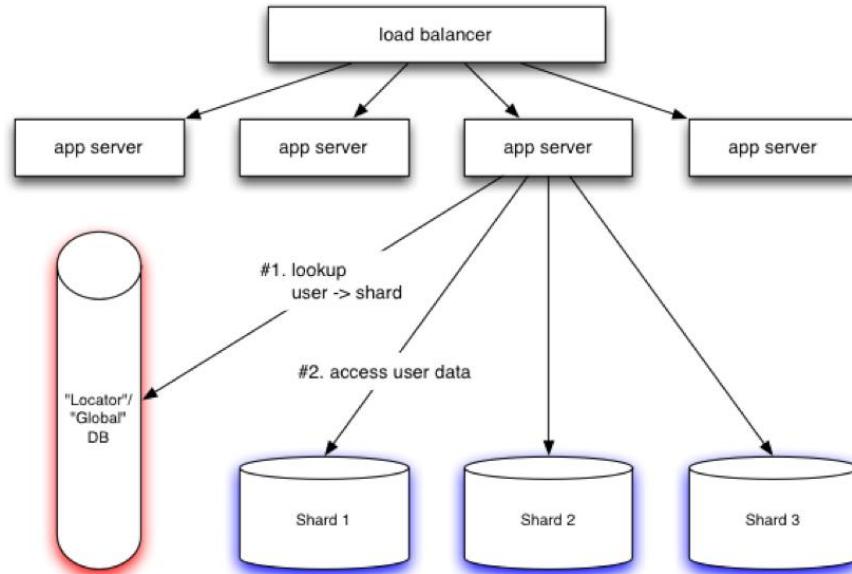
Sharding



Classic Sharding



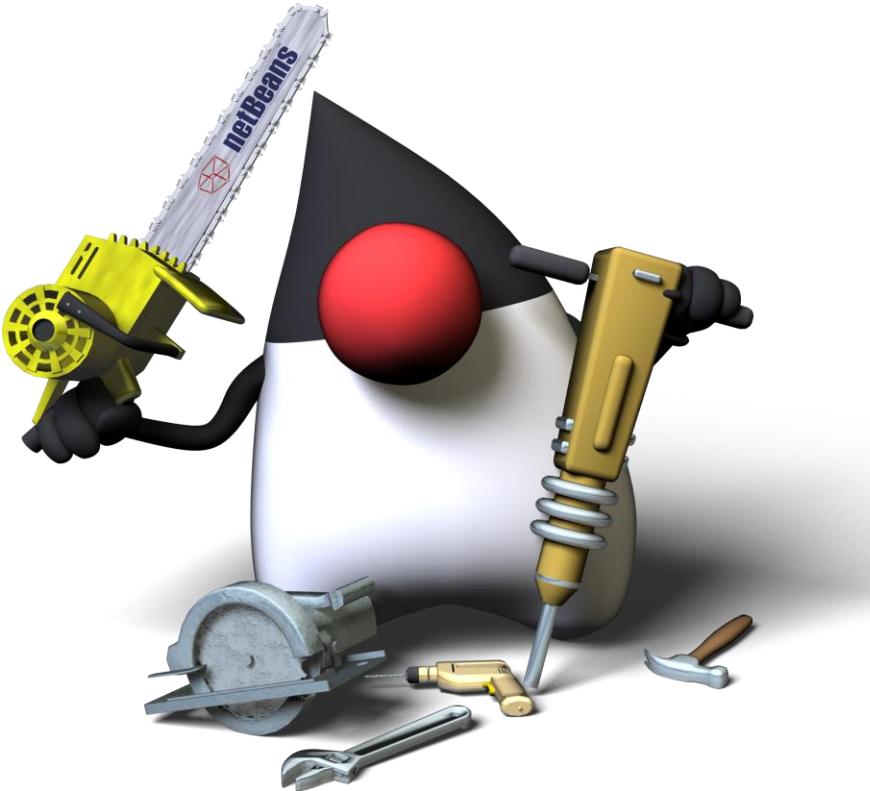
Sharding with Locator



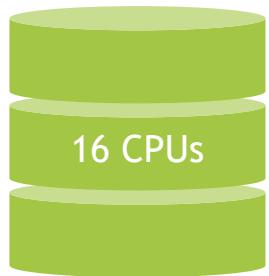
You'd measure performance!



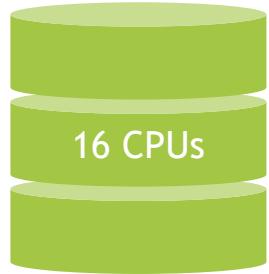
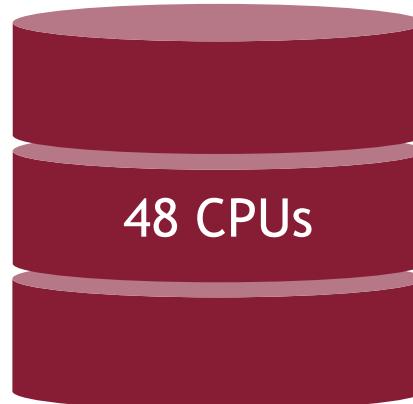
Case #3 : Supercomputer for President



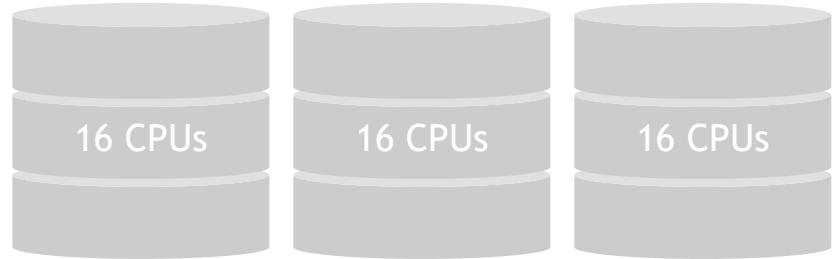
Motivation: Scale-up vs scale-out



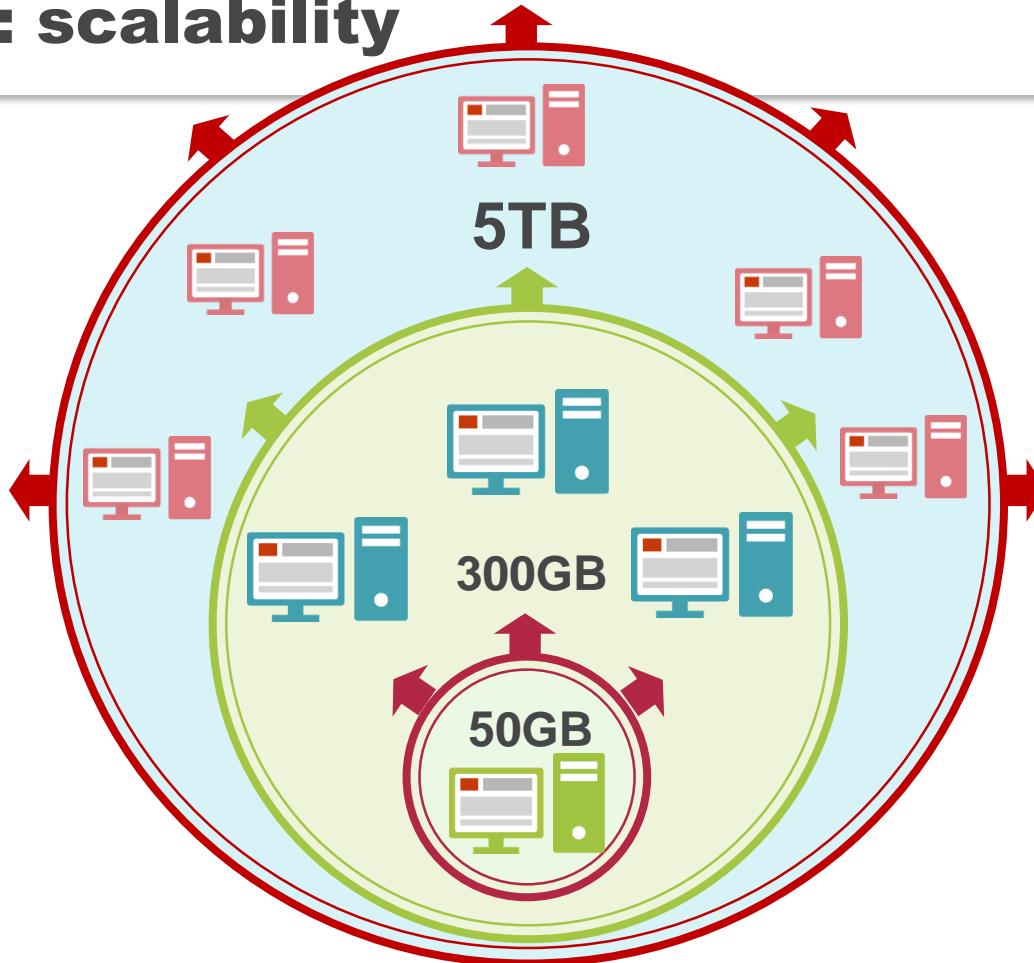
Scale - Up



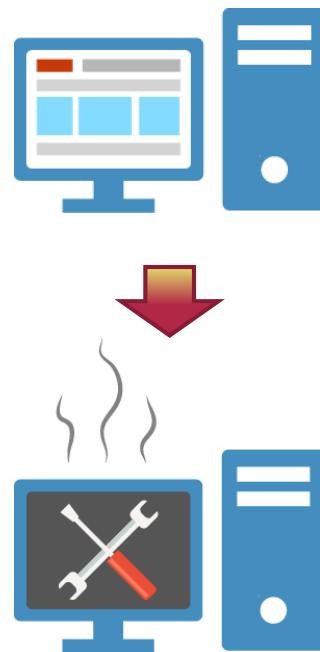
Scale - Out



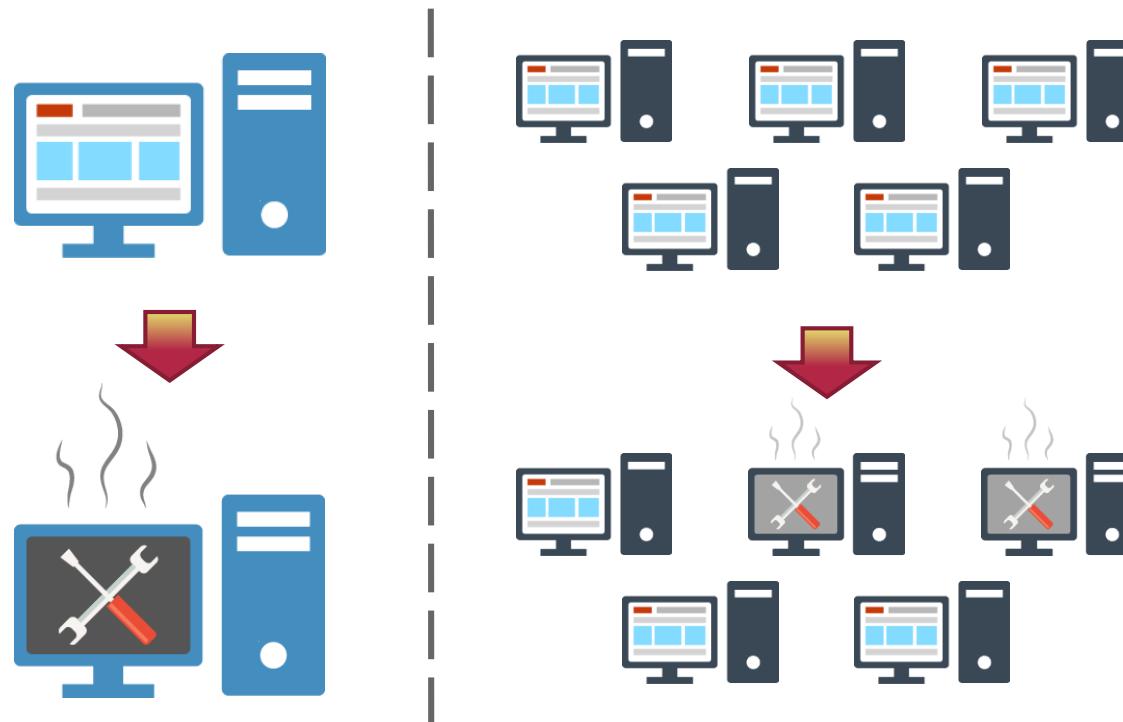
Motivation: scalability



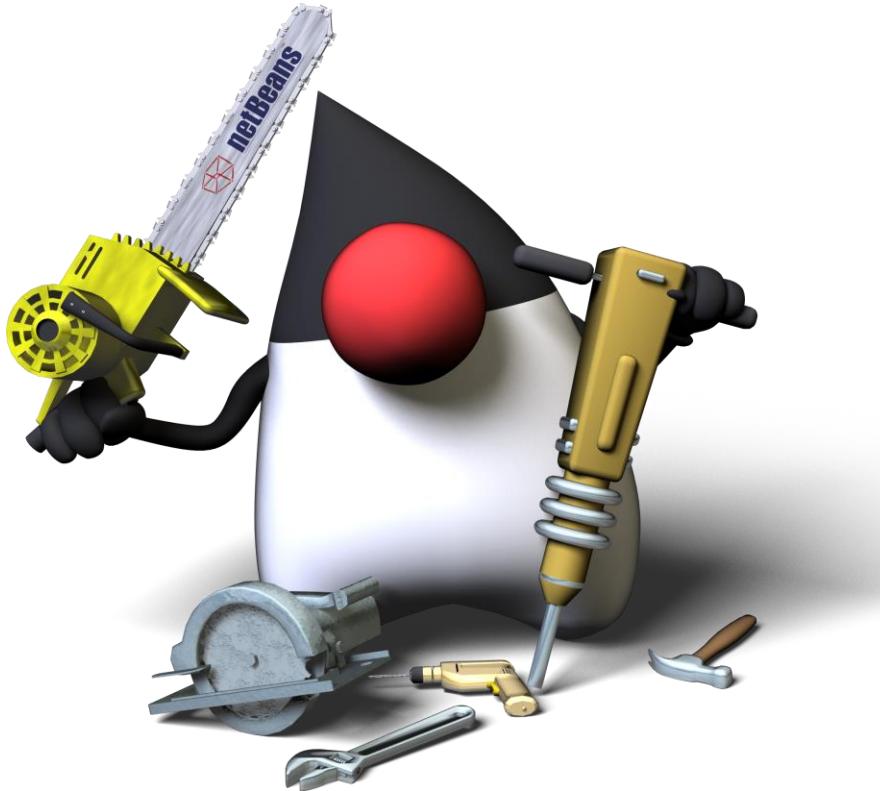
Motivation: Fault tolerance



Motivation: Fault tolerance

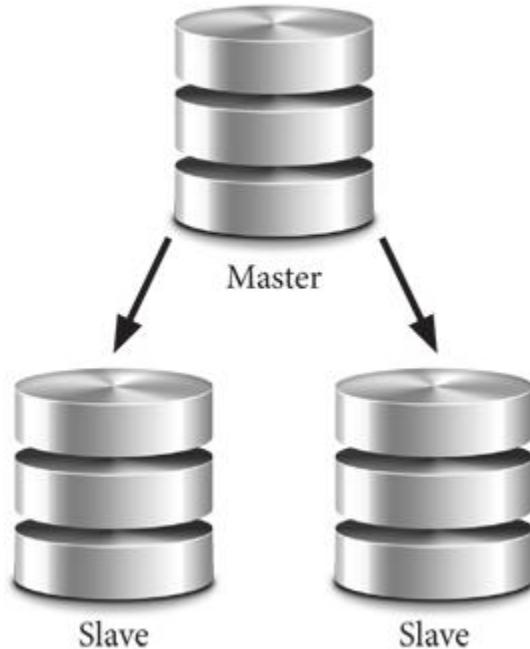


Case #4 : DDos attack against retail company

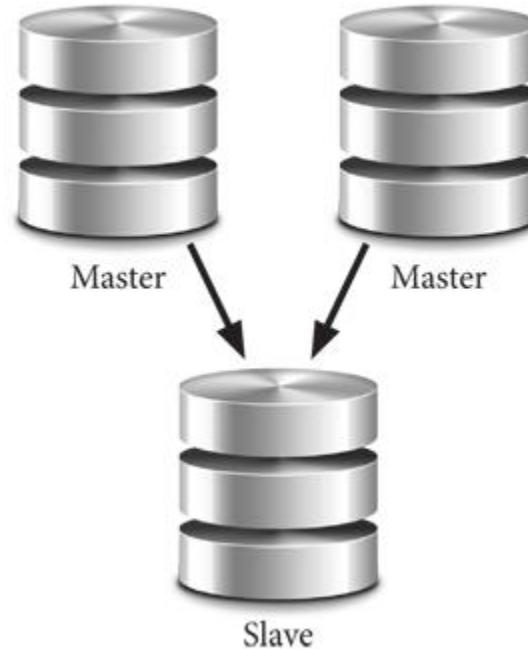


Replication

Replication



Multi-Source Replication



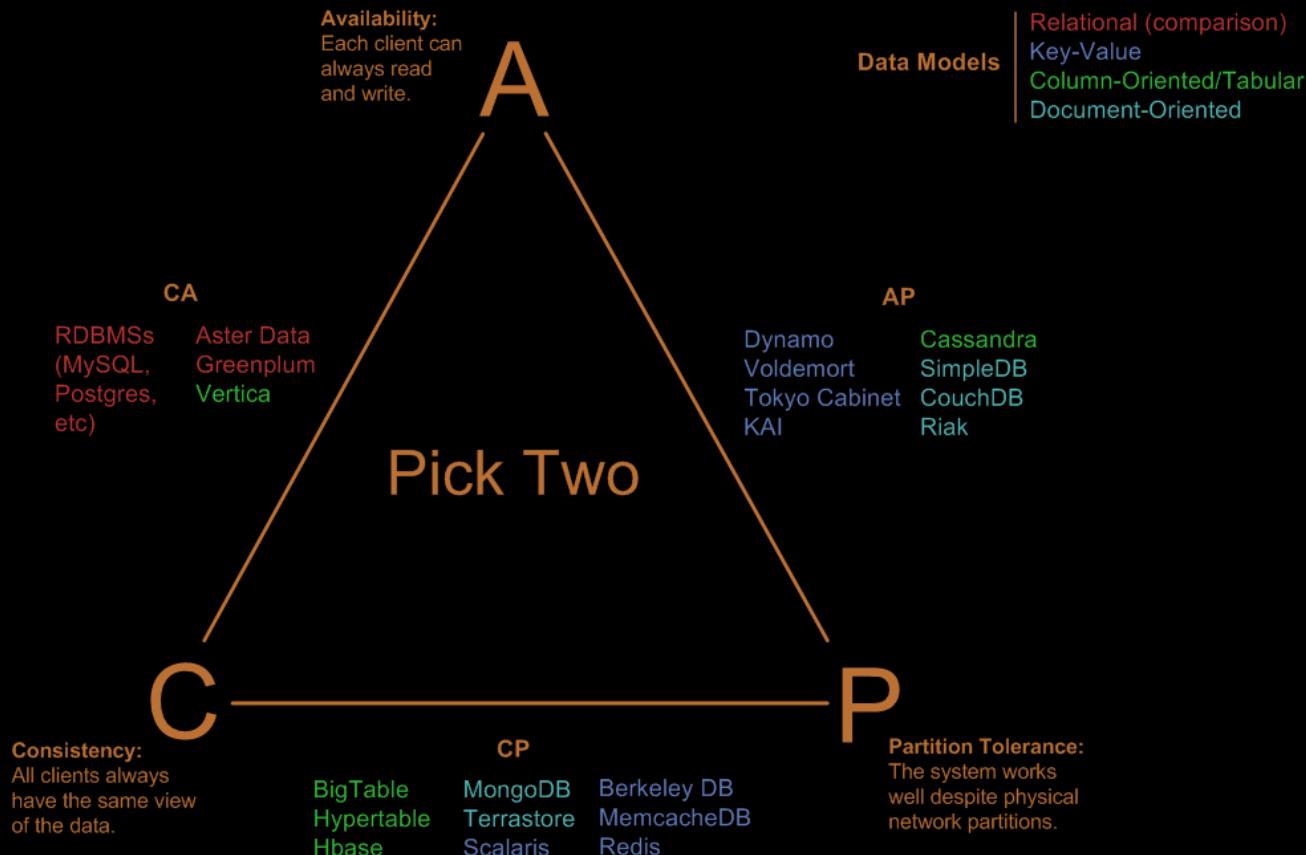
NOSQL

Replication



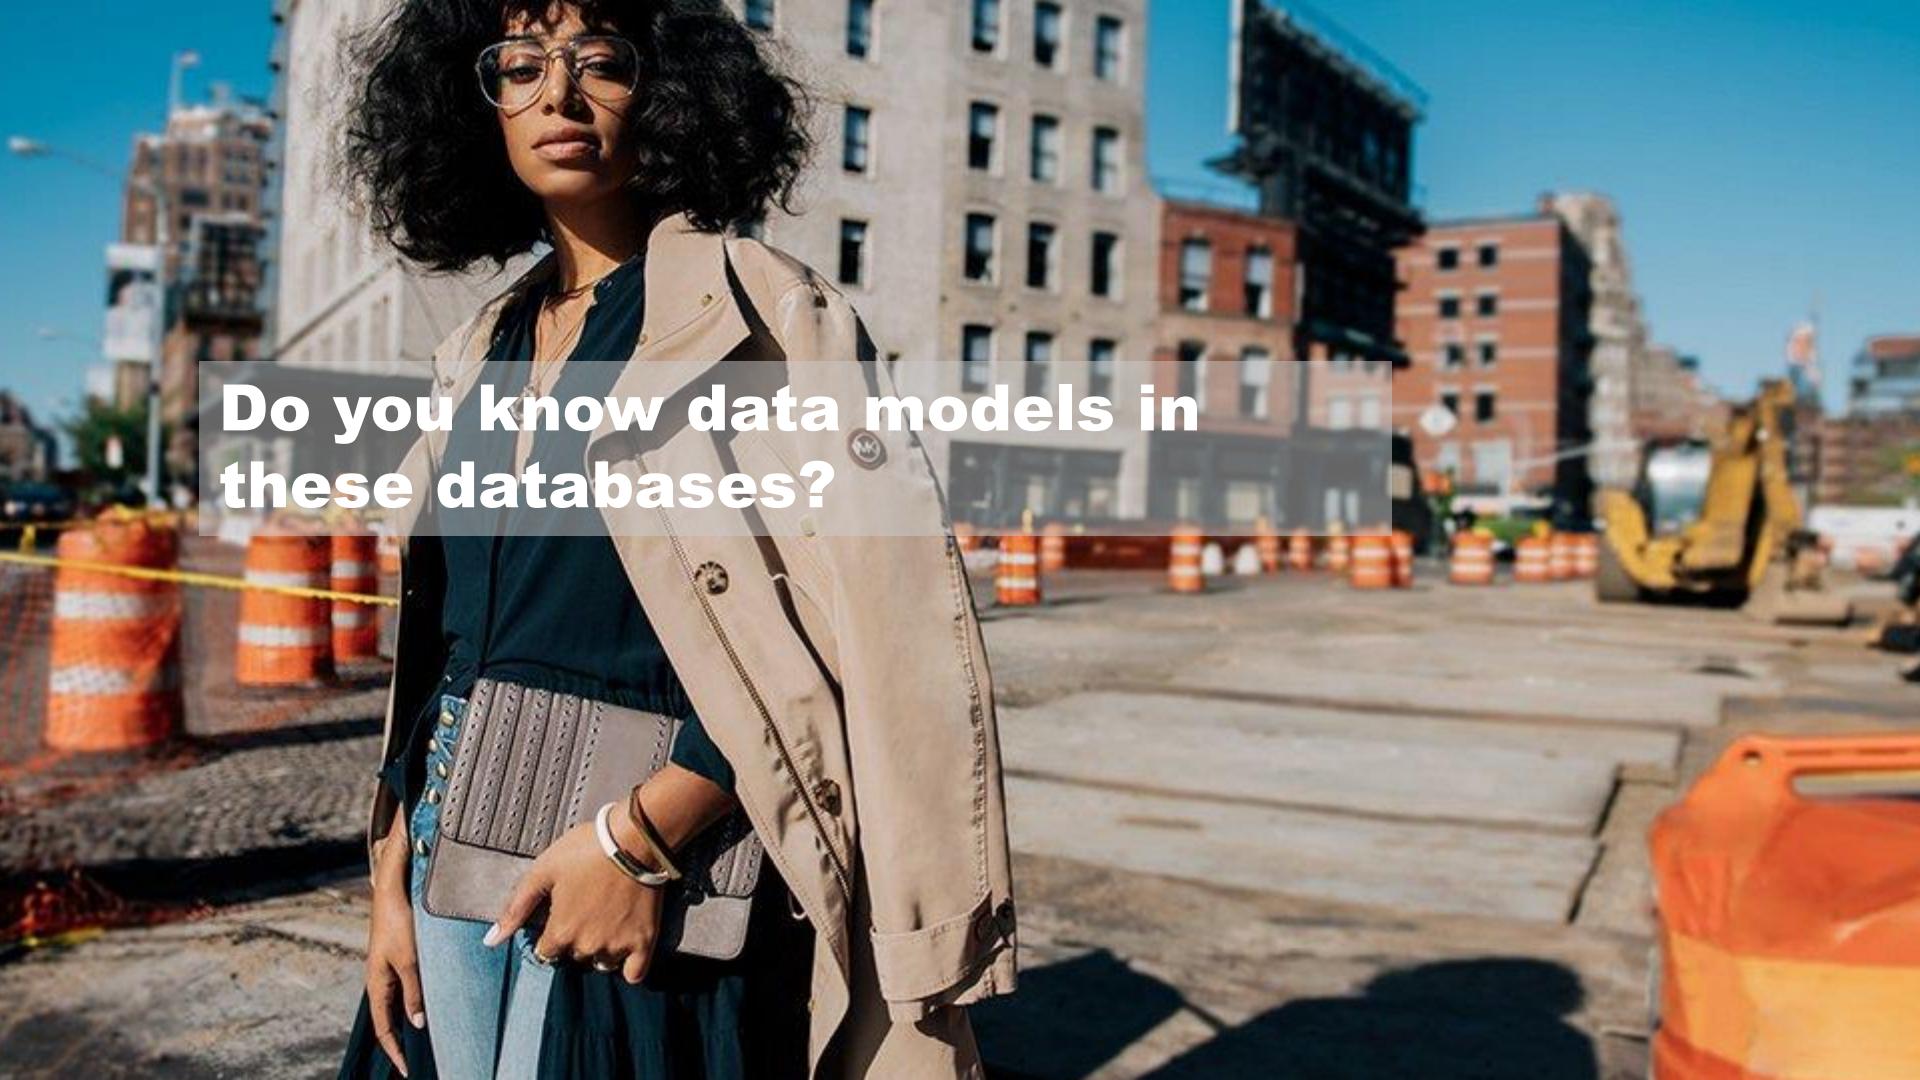
Visual Guide to NoSQL Systems

CAP Theorem



Databases

- Cassandra
- Hbase
- Neo4j
- Voldemort

A woman with dark curly hair and glasses, wearing a tan trench coat over a teal top and jeans, stands in the foreground of a city street under construction. She has her hands in her pockets and is looking towards the camera. In the background, there are several multi-story buildings, orange traffic cones, and a yellow excavator. A large white rectangular box covers the upper portion of the image, containing the text.

**Do you know data models in
these databases?**

Databases with types

- Cassandra (column family)
- Hbase (column family)
- Neo4j (graph)
- Voldemort (key-value)

Database party

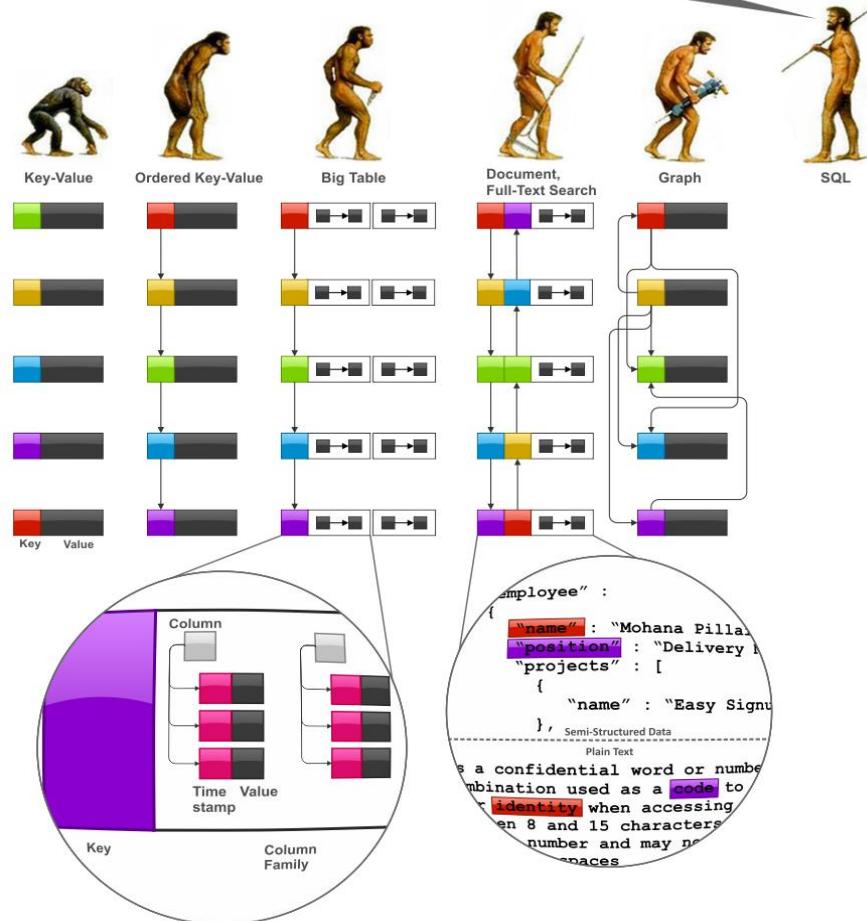


What's the problem with RDBMS's

- Caching
- Master/Slave
- Cluster
- Table Partitioning
- Sharding

Evolution

Stop following me, you fucking freaks!



Flowers in your garden

Data Model	Performance	Scalability	Flexibility	Complexity	Functionality
Key-value Stores	high	high	high	none	variable (none)

Flowers in your garden

Data Model	Performance	Scalability	Flexibility	Complexity	Functionality
Key-value Stores	high	high	high	none	variable (none)
Column Store	high	high	moderate	low	minimal

Flowers in your garden

Data Model	Performance	Scalability	Flexibility	Complexity	Functionality
Key-value Stores	high	high	high	none	variable (none)
Column Store	high	high	moderate	low	minimal
Document Store	high	variable (high)	high	low	variable (low)

Flowers in your garden

Data Model	Performance	Scalability	Flexibility	Complexity	Functionality
Key-value Stores	high	high	high	none	variable (none)
Column Store	high	high	moderate	low	minimal
Document Store	high	variable (high)	high	low	variable (low)
Graph Database	variable	variable	high	high	graph theory

Flowers in your garden

Data Model	Performance	Scalability	Flexibility	Complexity	Functionality
Key-value Stores	high	high	high	none	variable (none)
Column Store	high	high	moderate	low	minimal
Document Store	high	variable (high)	high	low	variable (low)
Graph Database	variable	variable	high	high	graph theory
Relational Database	variable	variable	low	moderate	relational algebra

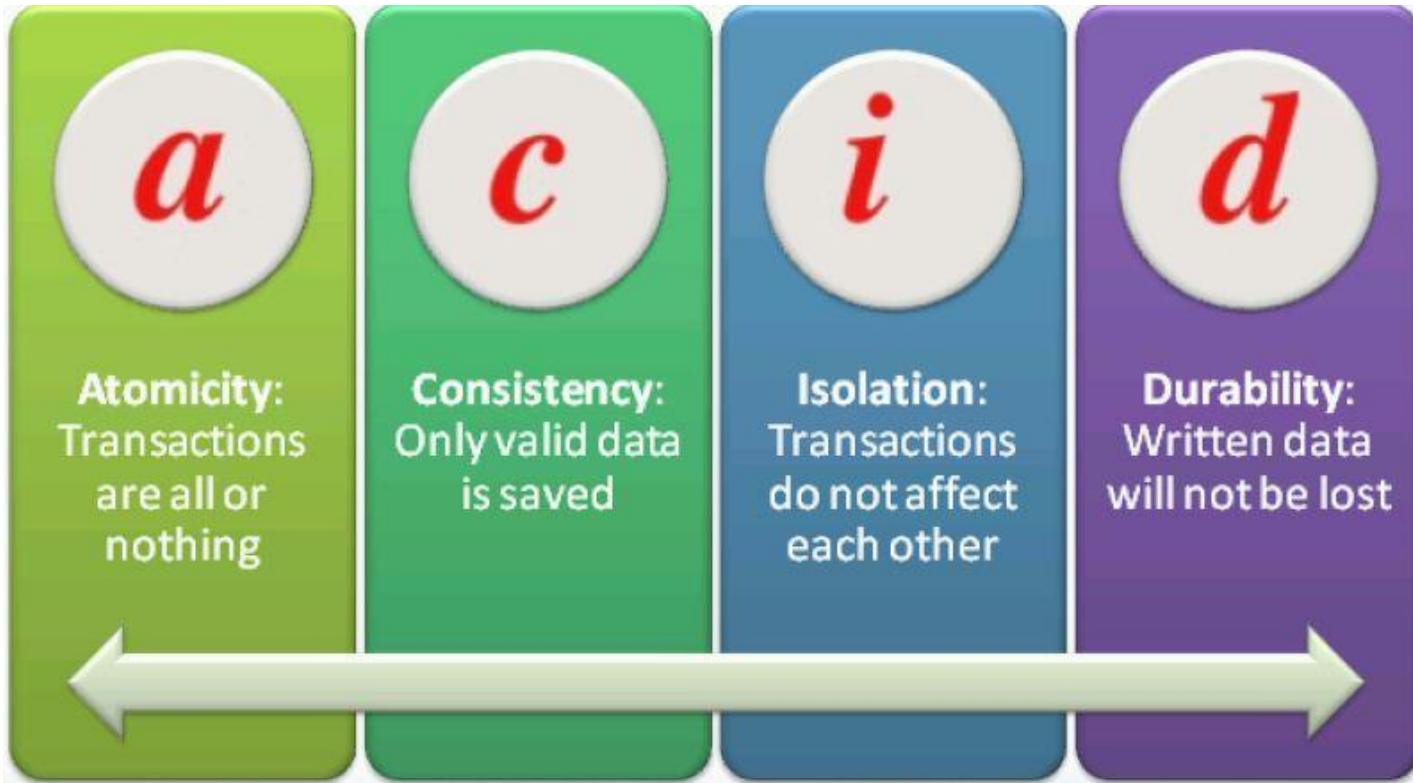
Gentle NoSQL



le-positive.ru

- Scalability
- Nodes and Data Centers
- Easy to add new server
- Specific data model
- Eventual consistency

ACID in SQL



Atomicity in NoSQL

- read-write-modify (CAS)
- key/row manipulation is atomic
- API for atomic operations
- bad support of transactions

- *basic availability* - all queries will be finished
- *soft state* - state can be changed without writing
- *eventual consistency*

Flowers in your garden

Database	Data model	Query API	Data storage system
Cassandra	Column Family	Thrift	Memtable/SSTable

Flowers in your garden

Database	Data model	Query API	Data storage system
Cassandra	Column Family	Thrift	Memtable/SSTable
CouchDB	Documents	Map/Reduce	Append-only-B-tree

Flowers in your garden

Database	Data model	Query API	Data storage system
Cassandra	Column Family	Thrift	Memtable/SSTable
CouchDB	Documents	Map/Reduce	Append-only-B-tree
Hbase	Column Family	Thrift, REST	Memtable/SSTable on HDFS

Flowers in your garden

Database	Data model	Query API	Data storage system
Cassandra	Column Family	Thrift	Memtable/SSTable
CouchDB	Documents	Map/Reduce	Append-only-B-tree
Hbase	Column Family	Thrift, REST	Memtable/SSTable on HDFS
MongoDB	Documents	Cursor	B-tree

Flowers in your garden

Database	Data model	Query API	Data storage system
Cassandra	Column Family	Thrift	Memtable/SSTable
CouchDB	Documents	Map/Reduce	Append-only-B-tree
Hbase	Column Family	Thrift, REST	Memtable/SSTable on HDFS
MongoDB	Documents	Cursor	B-tree
Neo4j	Edges/Verticies	Graph	On-disk linked lists

Flowers in your garden

Database	Data model	Query API	Data storage system
Cassandra	Column Family	Thrift	Memtable/SSTable
CouchDB	Documents	Map/Reduce	Append-only-B-tree
Hbase	Column Family	Thrift, REST	Memtable/SSTable on HDFS
MongoDB	Documents	Cursor	B-tree
Neo4j	Edges/Verticies	Graph	On-disk linked lists
Riak	Key/Value	Nested hashes, REST	Hash



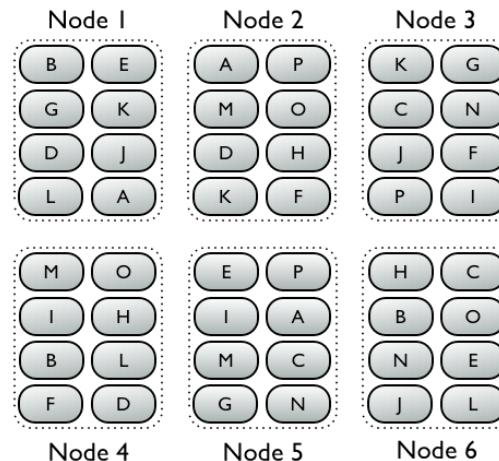
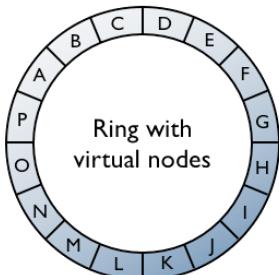
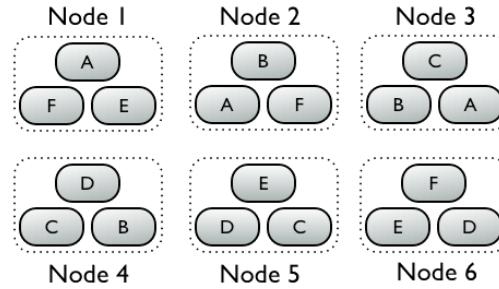
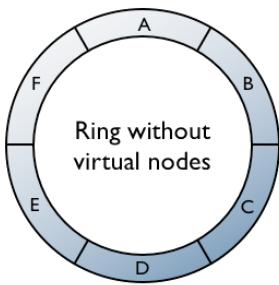
What about queries?

Mongo



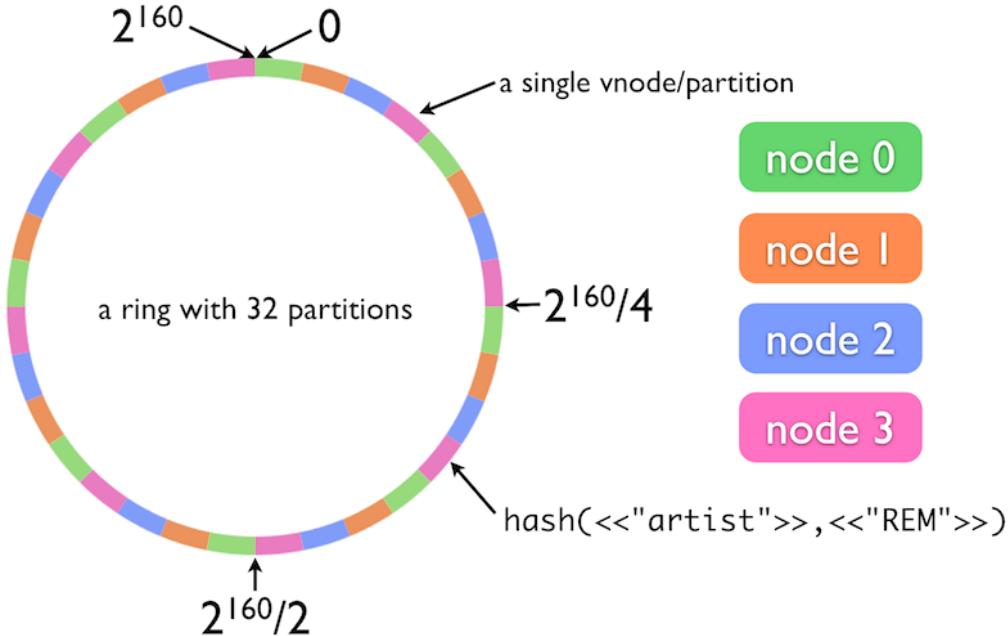
- 16 mb
- JavaScript at the bottom
- 2d, 3d, B-tree indexes
- 3.0 version is very hot
- integration with Kafka

Cassandra

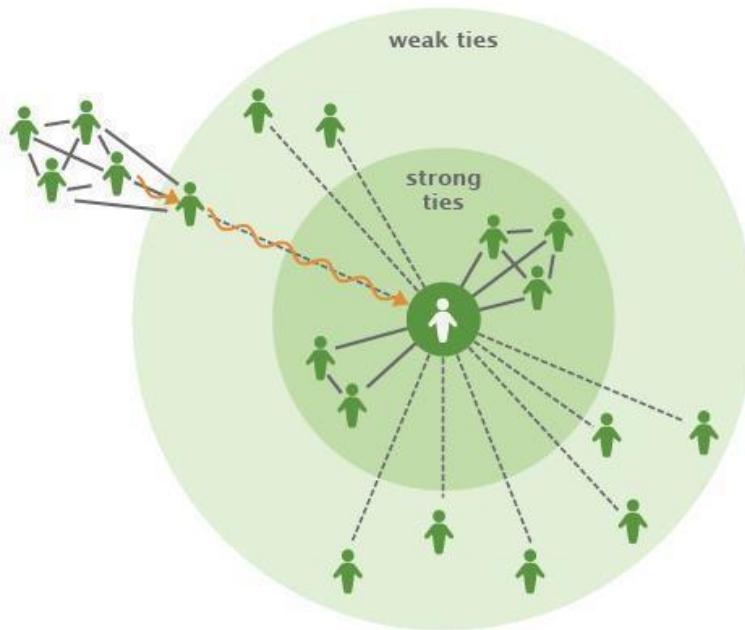


- CQL
- 2 billions columns in row
- No ACID, of course
- Can spend all your RAM
- JVM - based

Riak



- Links to another keys
- REST
- Consistency level in each query
- Ring of nodes



- Vertices, edges
- ACID
- REST API + Cypher
- 2d index
- Not so good for distributed data

A woman with dark curly hair and glasses, wearing a tan trench coat over a teal top and jeans, stands in the foreground of a city street under construction. She is looking towards the camera. The background shows several multi-story buildings, orange and white traffic cones, and a yellow excavator. A large white rectangular box covers the middle portion of the image, containing the text.

And what should we choose?

Network Rule

Can your data be presented as network or graph?

If yes -> Neo4j

If no -> continue

BigData Rule

Do you have TB or PB of data?

If yes -> NoSQL

If no -> Maybe you should stay with SQL?

Easy Rule

Do you need in more complex operation than read/write by key?

If yes -> continue

If no -> Riak, Memcached

Hierarchy Rule

Do you have nested data? Is R >> W in your system?

If yes -> MongoDB, CouchDB

If no -> continue

Hadoop Rule

Is Hadoop integration required? Can your data be presented as flat table?

If yes -> Hbase

If no -> continue

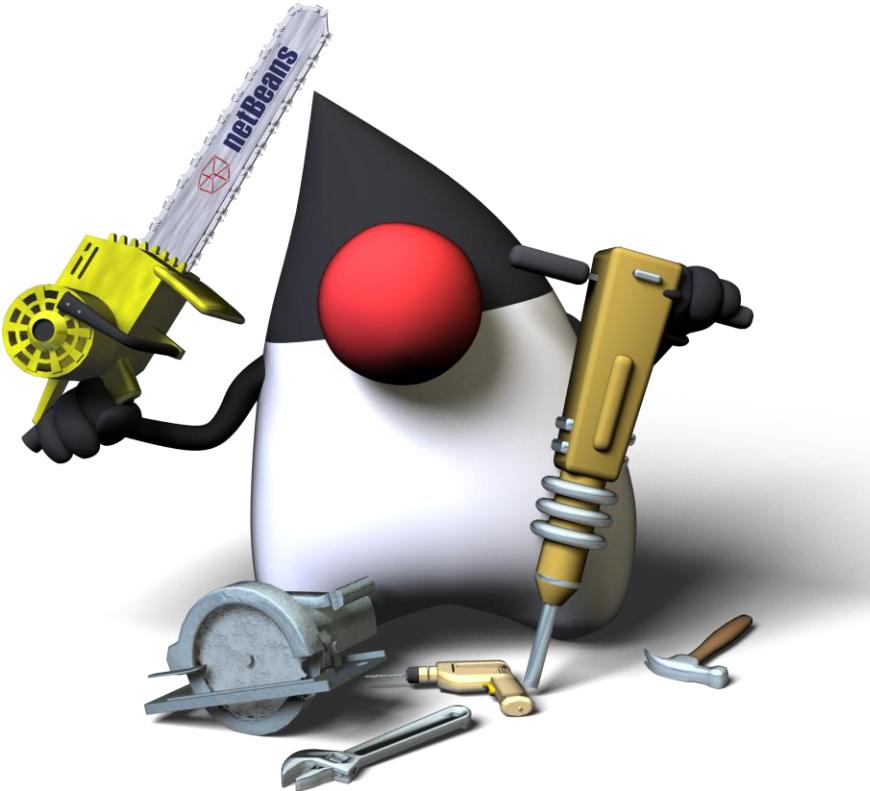
Availability Rule

Do you need in high availability? Are you agree with eventual consistency?

If yes -> Cassandra

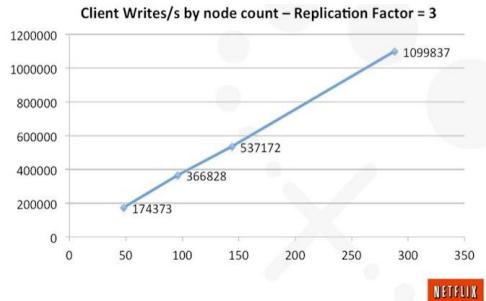
If no -> continue

Case #5 : Go back to PostgreSQL

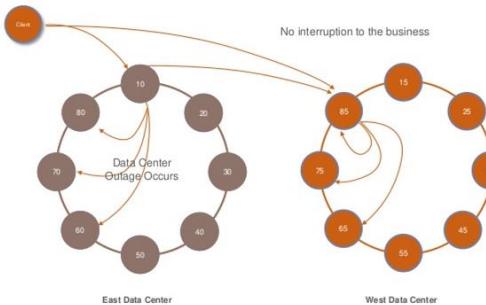


CASSANDRA

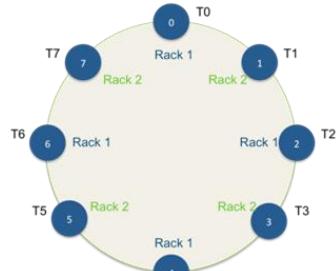
Cassandra Features



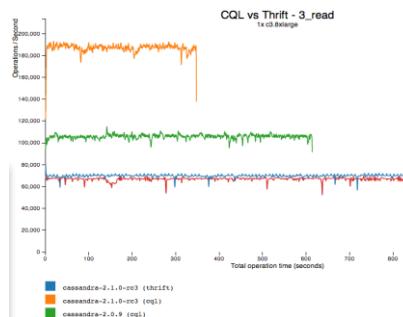
LINEAR SCALABILITY



MULTI DATA CENTER SUPPORT



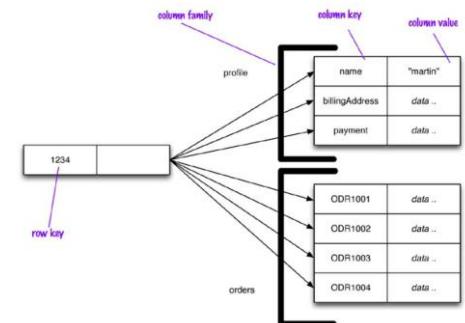
MASTER-LESS ARCHITECTURE



MASSIVE WRITE AND FAST RESPONSE

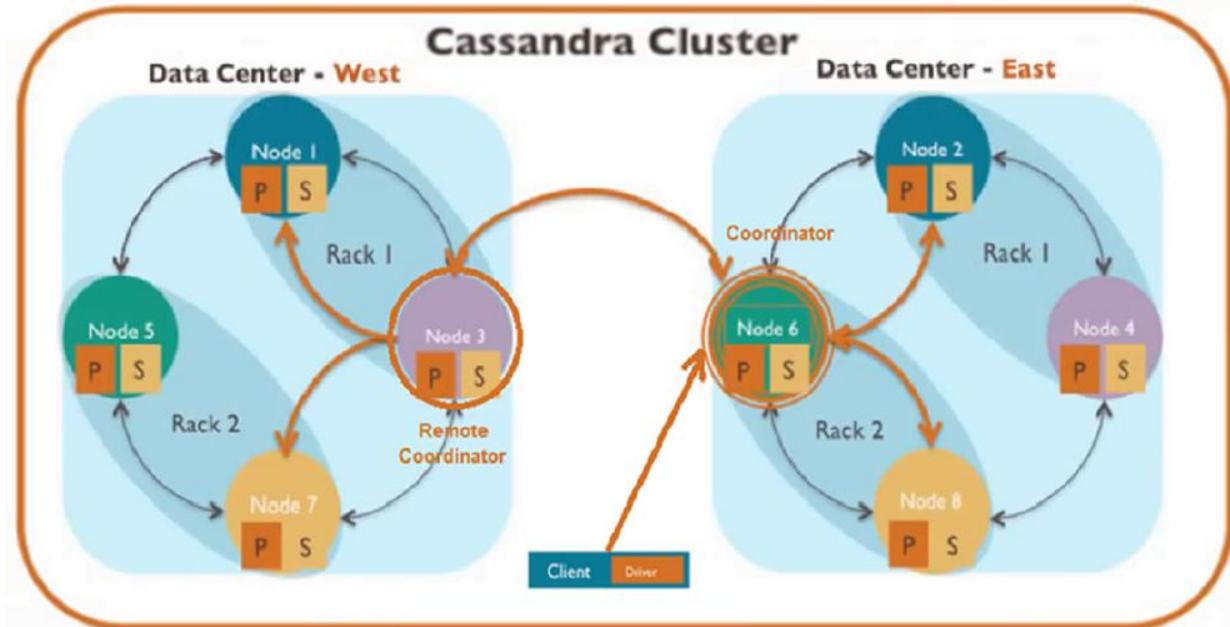
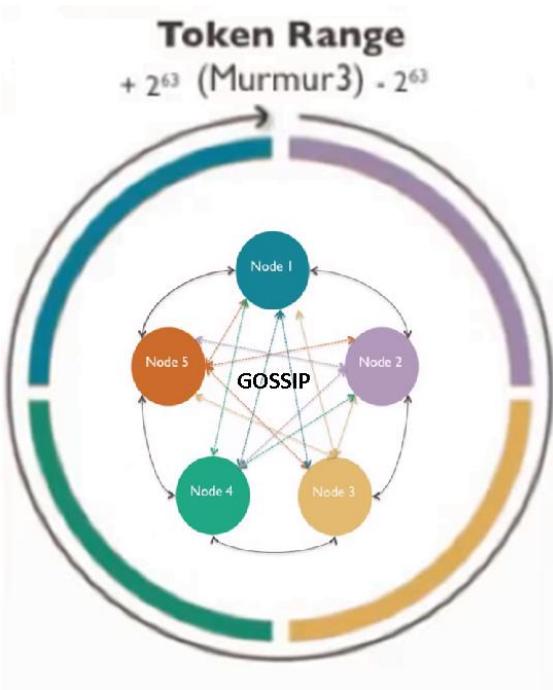
ALL	QUORUM	ONE
Strong Consistency	Quorum == (nodes / 2) + 1	High Availability / Eventual Consistency

TUNABLE CONSISTENCY



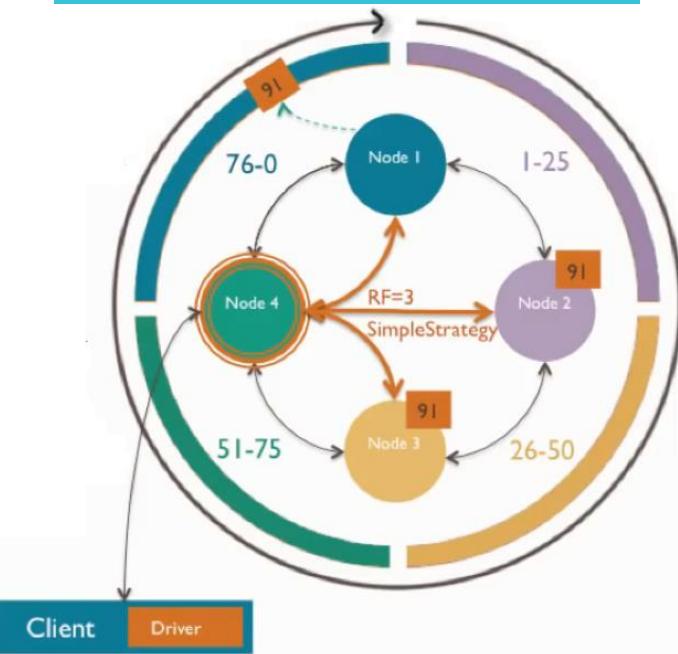
FLEXIBLE DATA MODEL

Cassandra “Ring” Architecture



Cassandra Replication

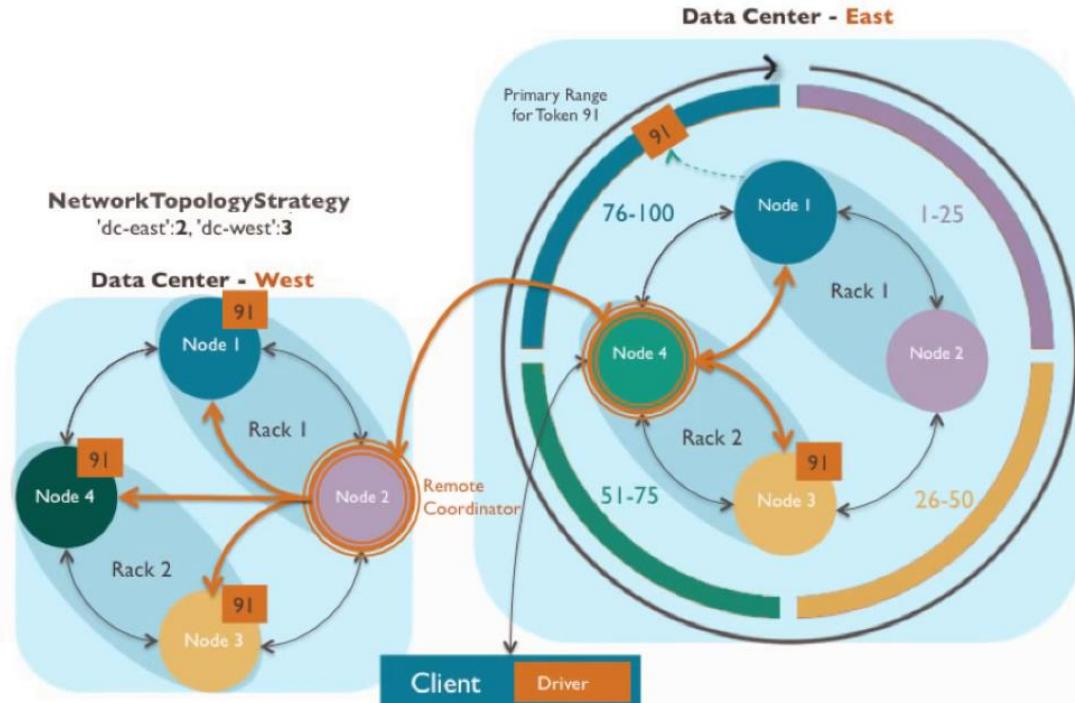
SimpleStrategy



NetworkTopologyStrategy
'dc-east':2, 'dc-west':3

Data Center - West

Data Center - East



NetworkTopologyStrategy

Cassandra Tunable Consistency

Write Consistency

ALL

EACH_QUORUM

QUORUM

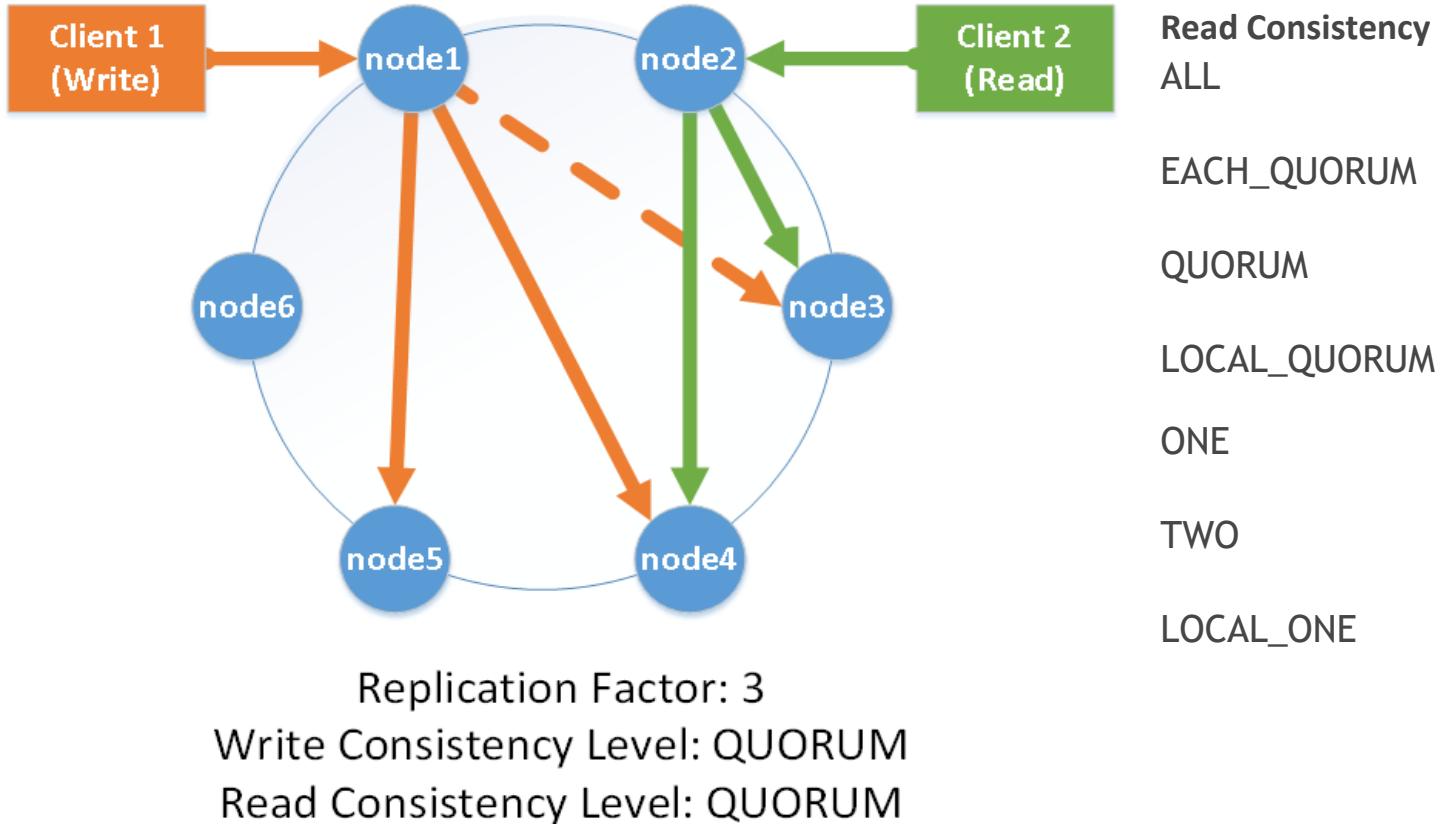
LOCAL_QUORUM

ONE

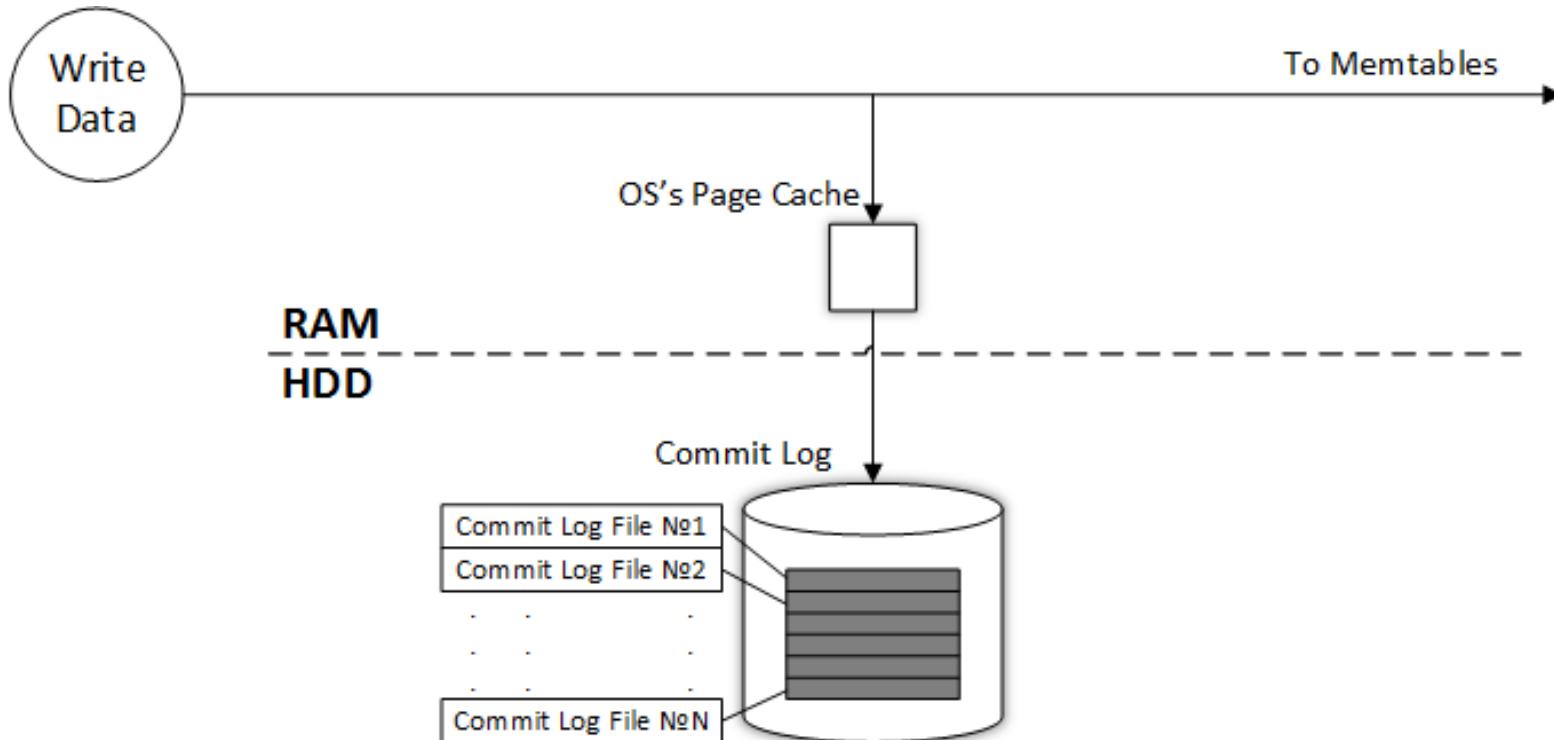
TWO

LOCAL_ONE

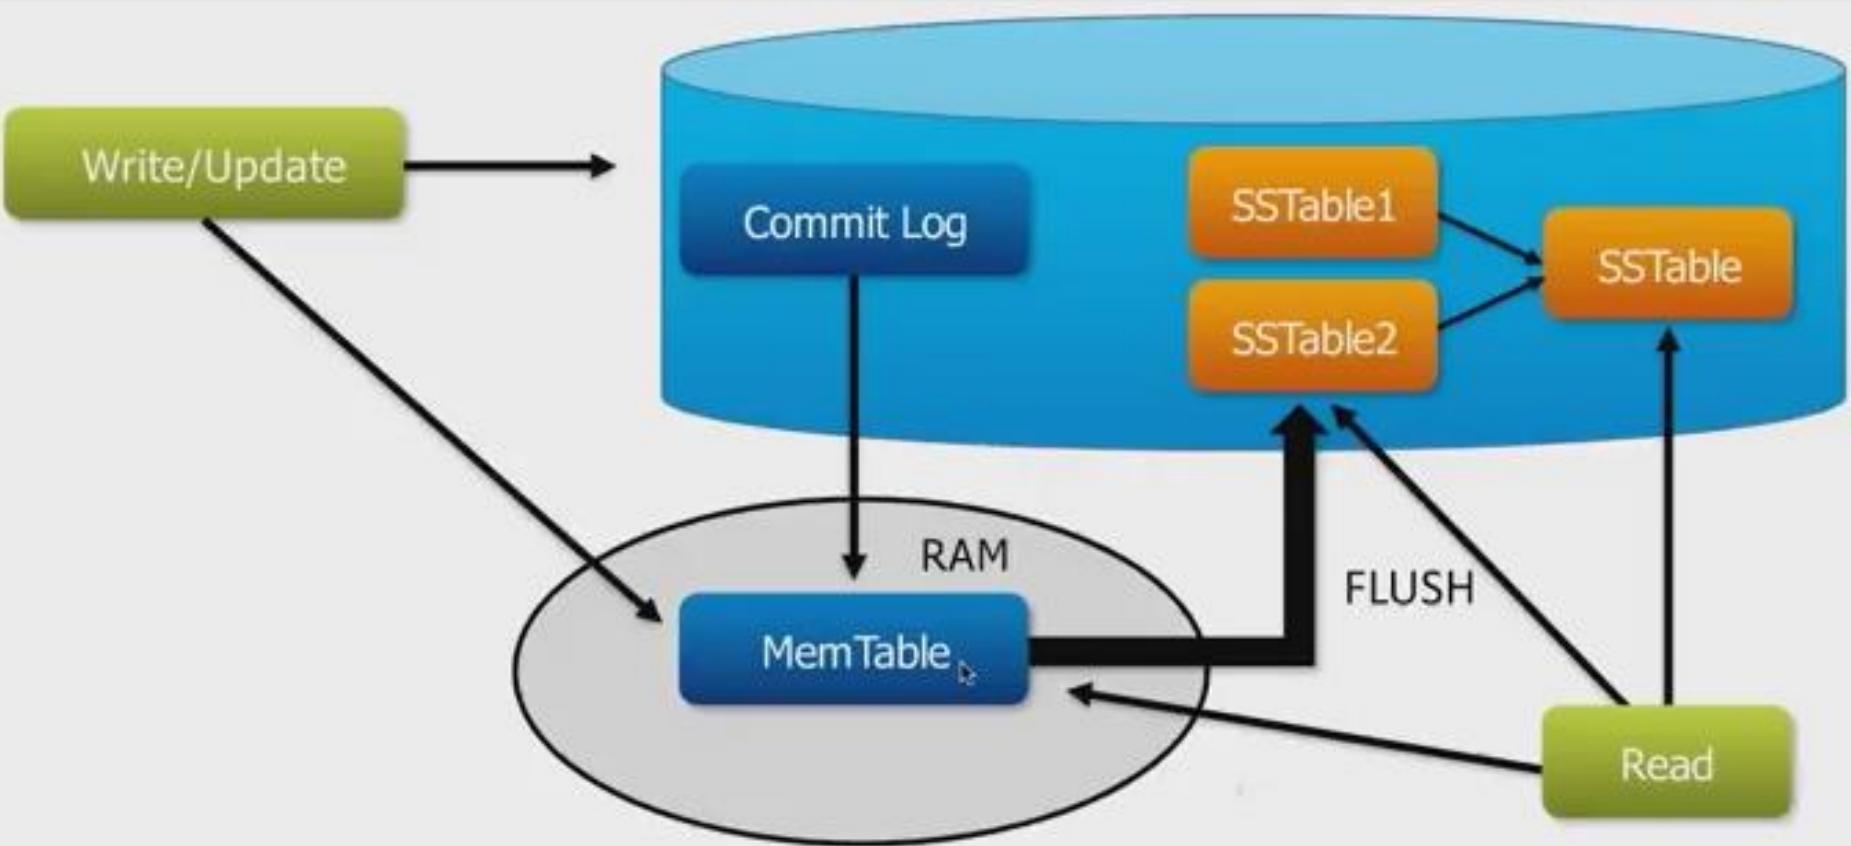
ANY



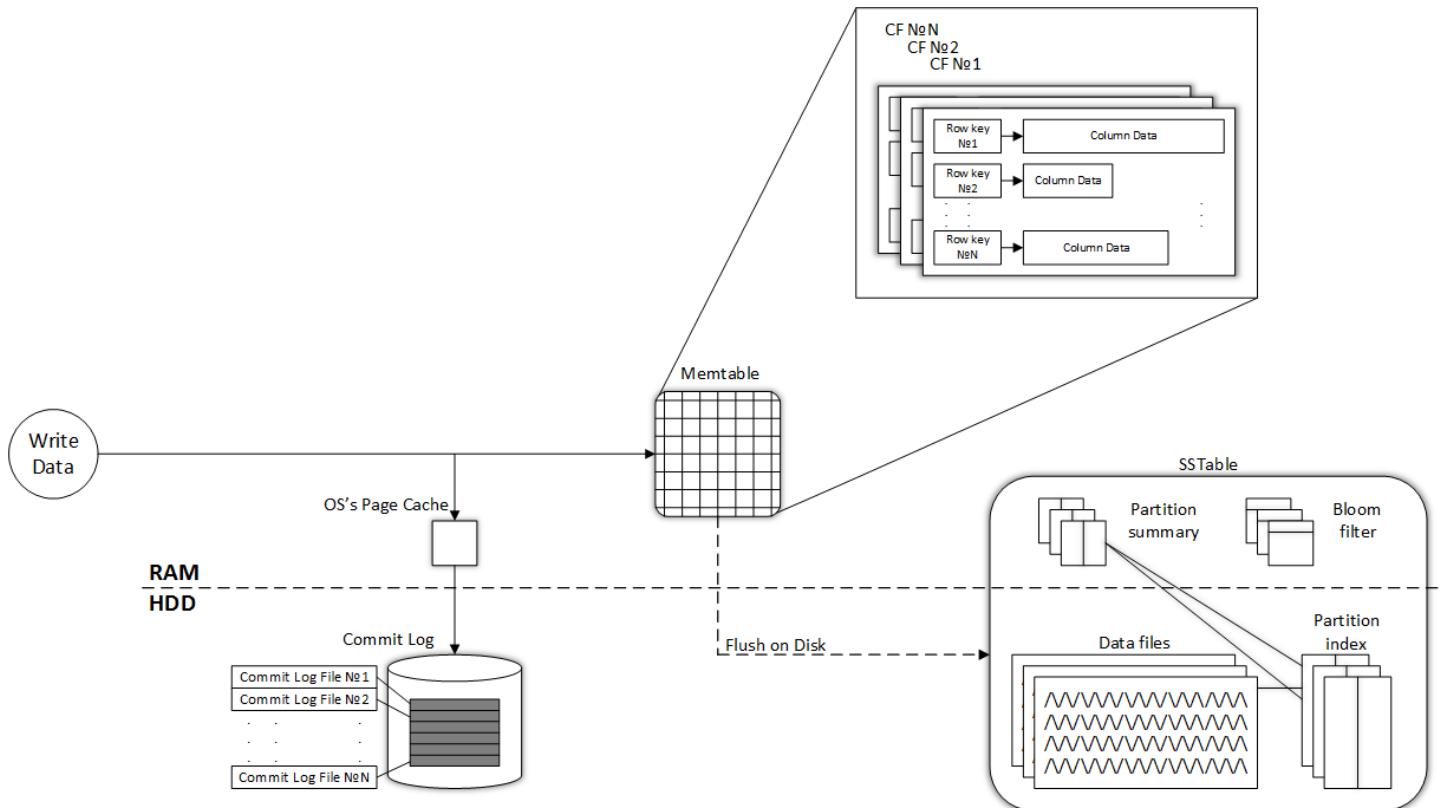
Cassandra write path (COMMITLOG)



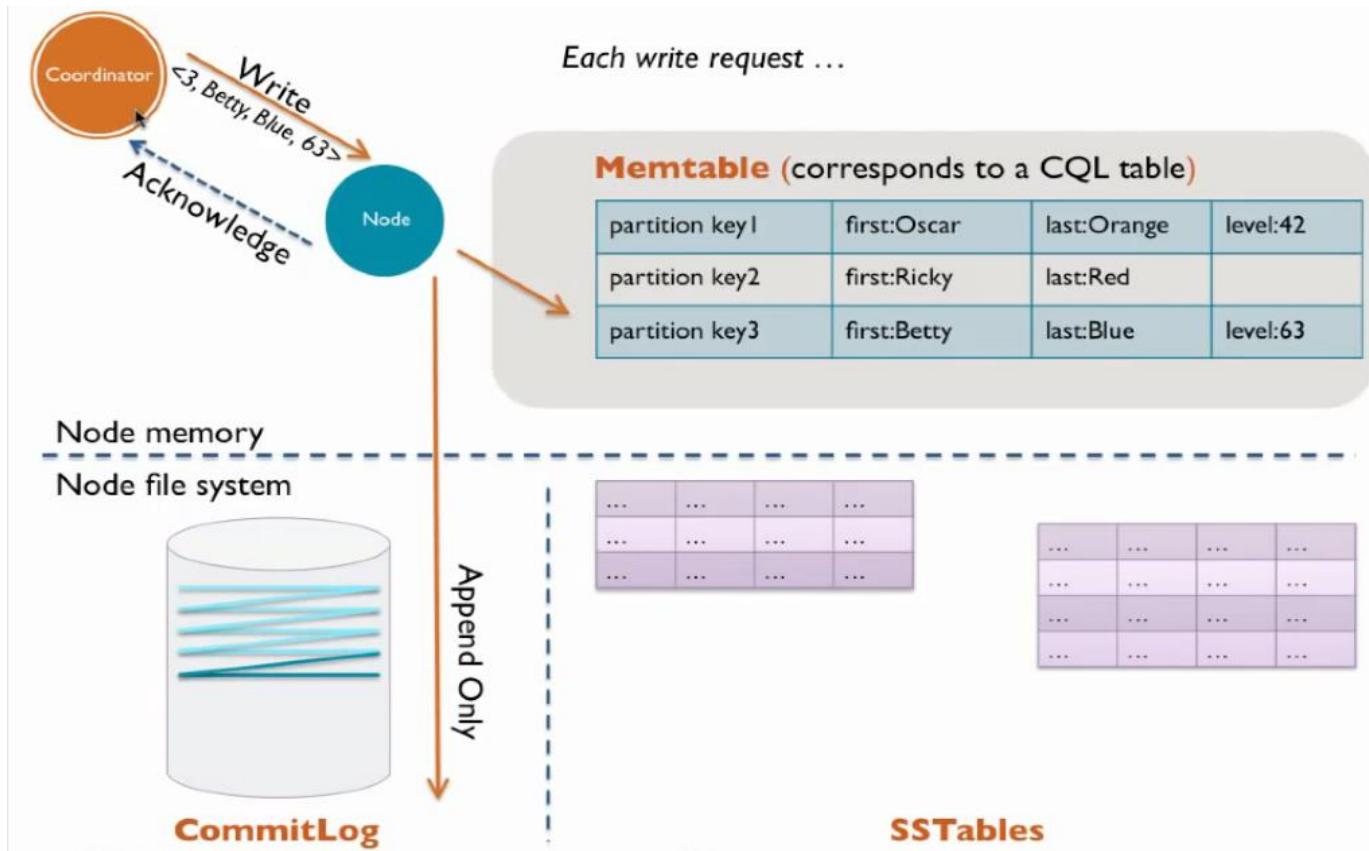
Cassandra write path (MemTable)



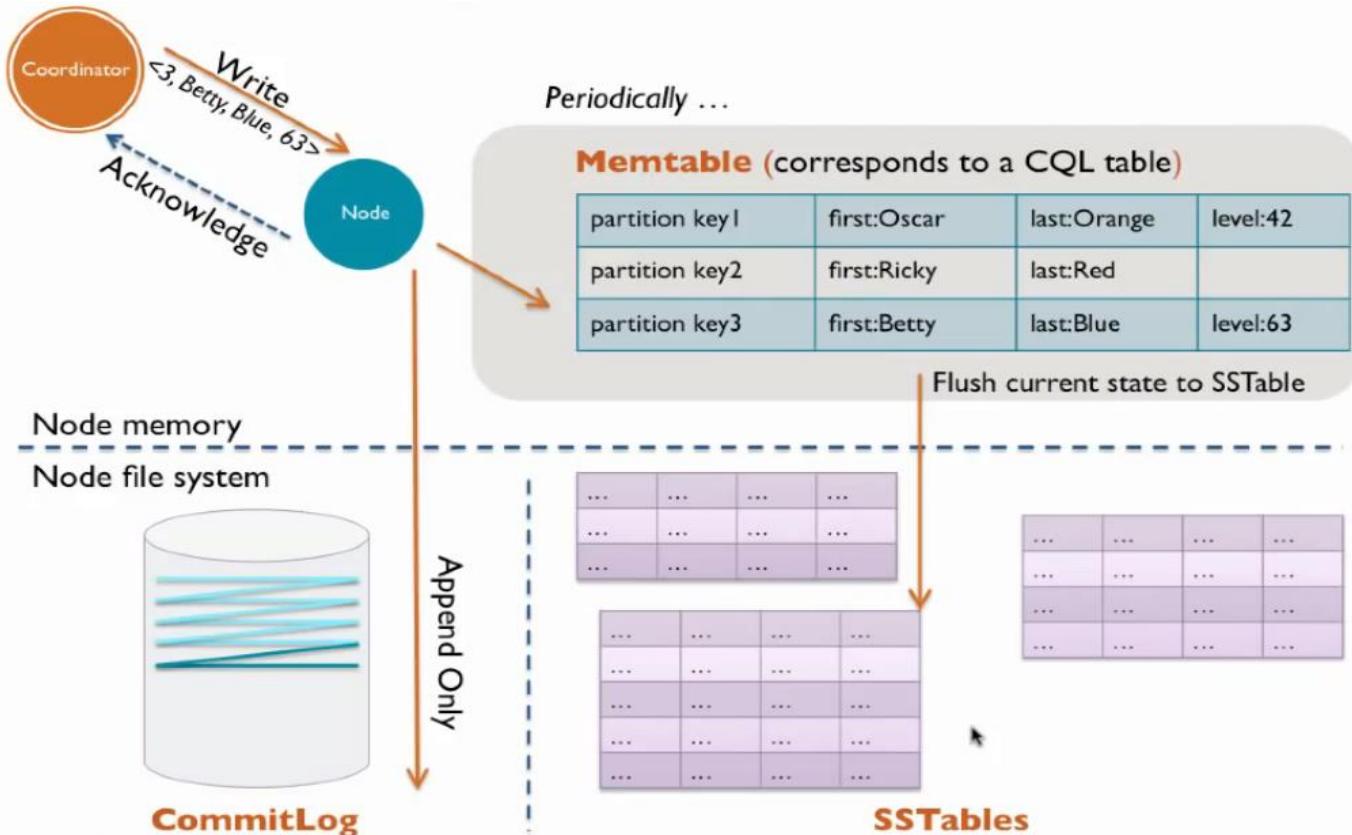
Cassandra write path (SSTables)



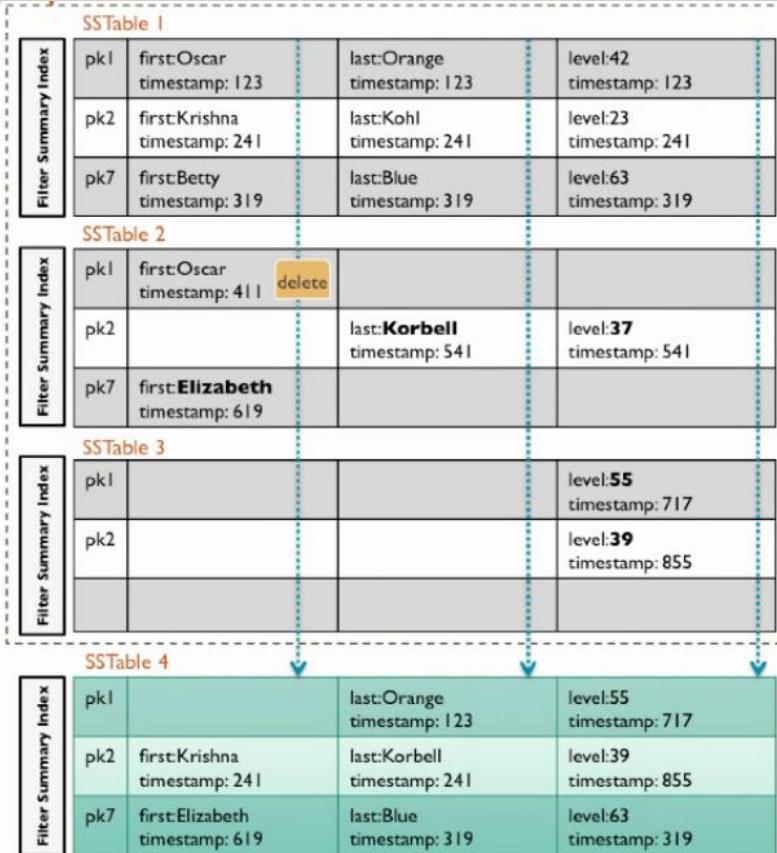
Cassandra write path



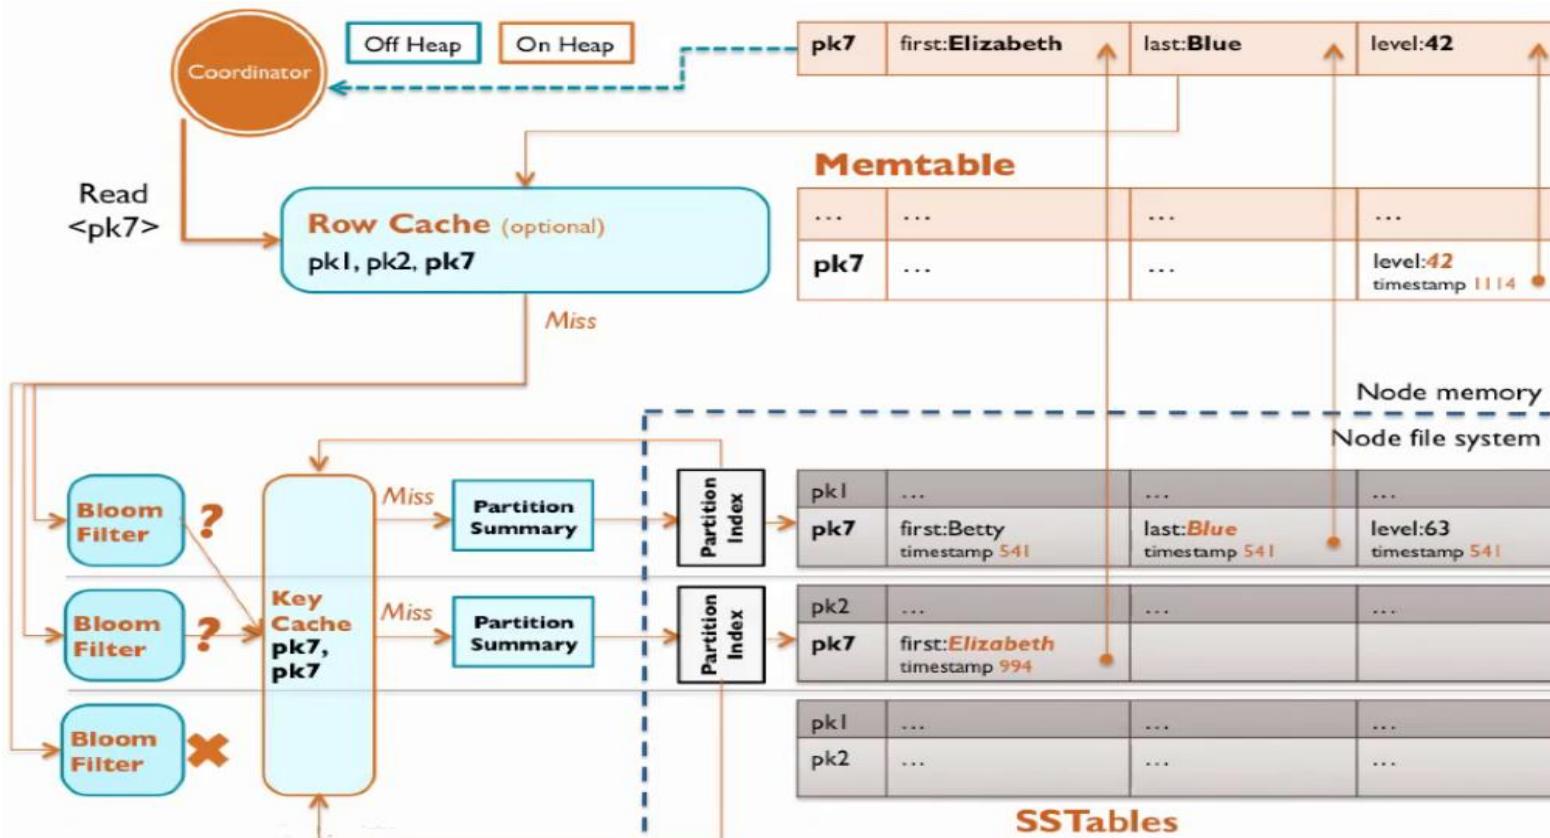
Cassandra write path (FLUSH)



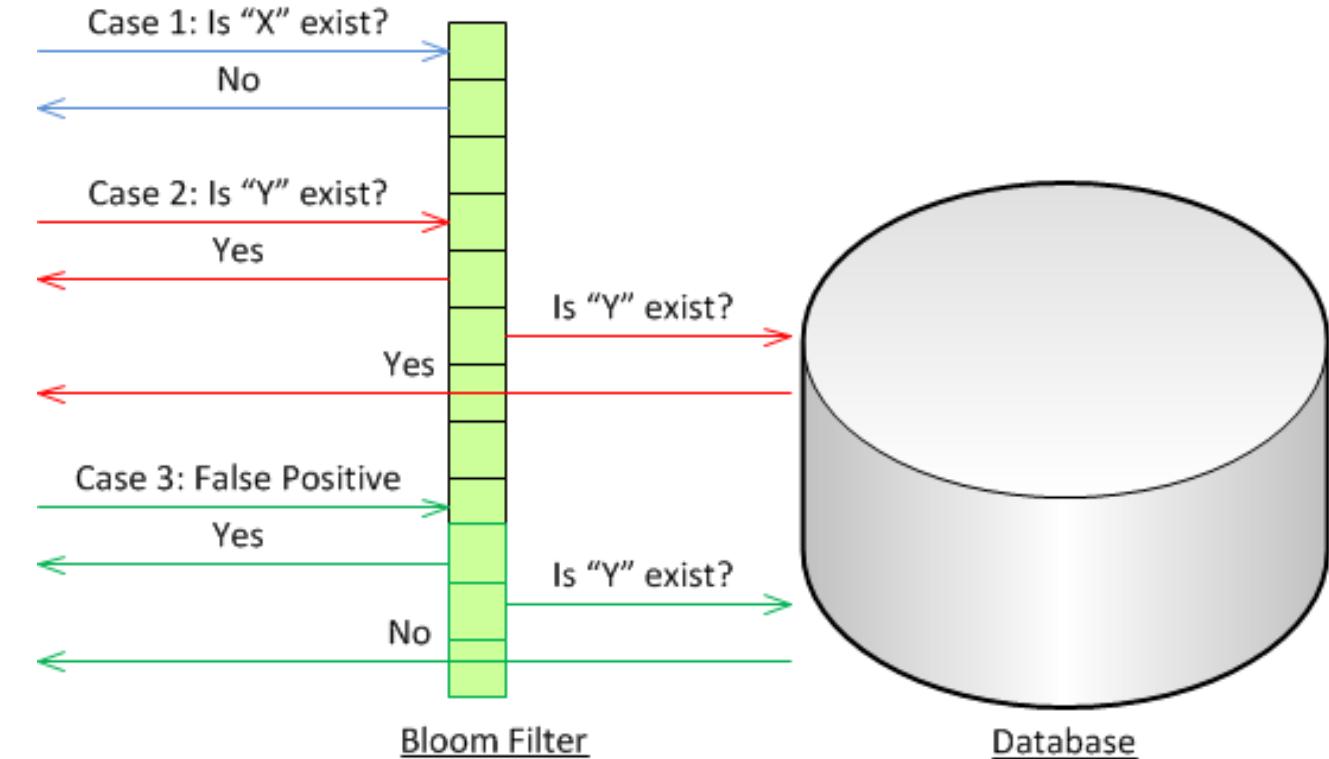
Cassandra write path (COMPACTIOn)



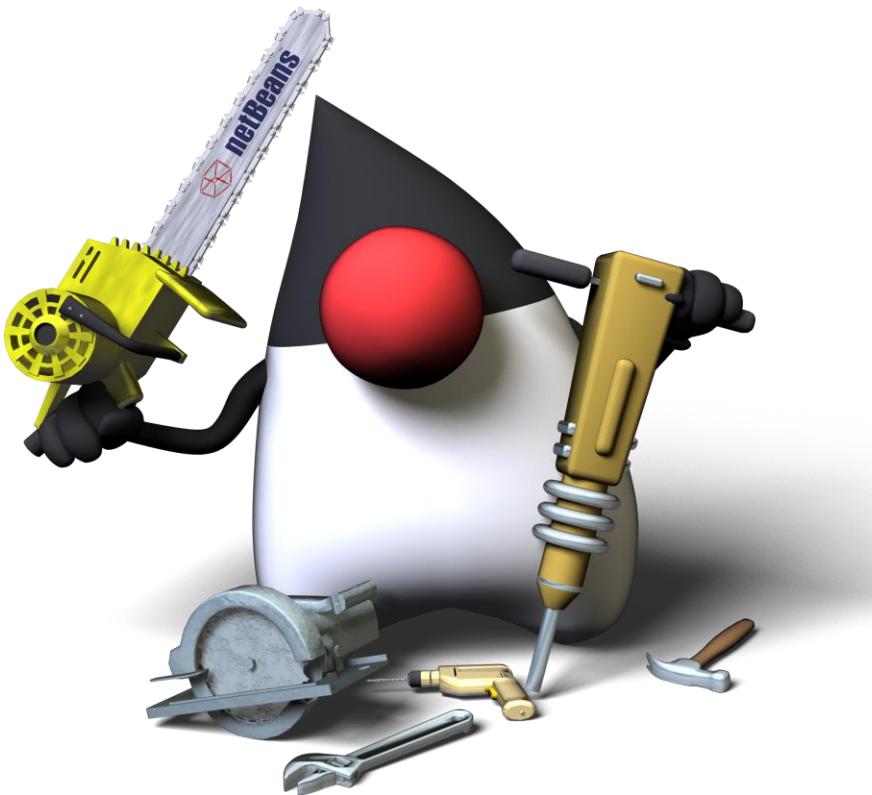
Cassandra READ path



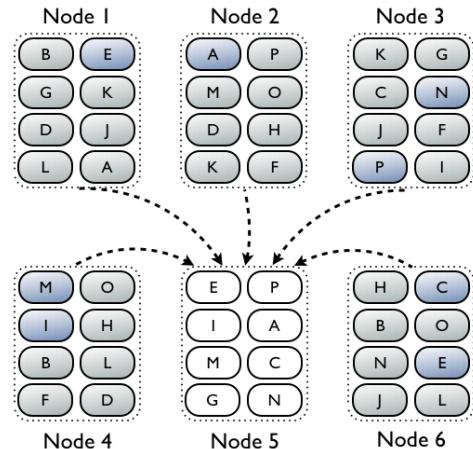
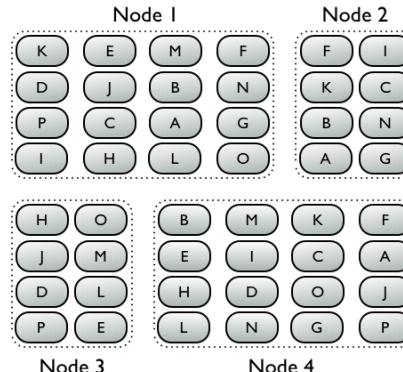
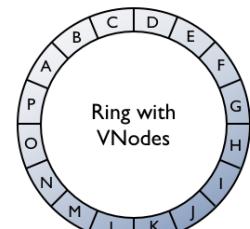
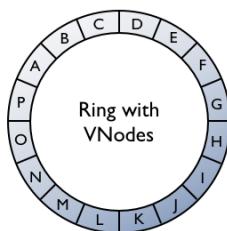
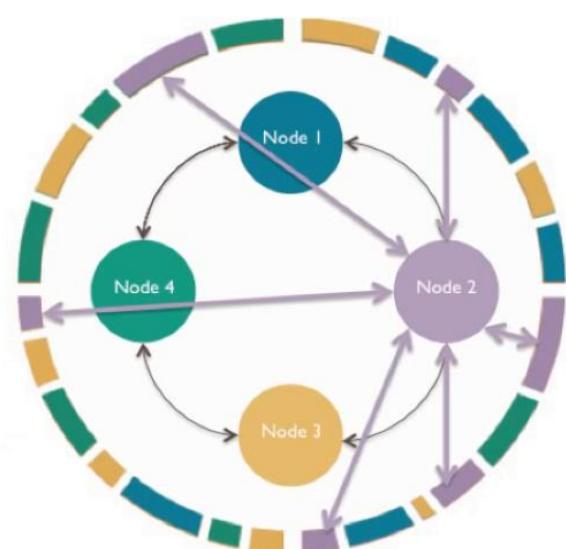
Bloom Filter



Case #6 : Strong Wall with Bloom Filters



Cassandra Virtual nodes



DATA MODEL

Data Model: Map<Map<Key, Value>>

Invoice
001

Buyer	X_1232_Name	X_1232_Price	X_4475_Name	X_4475_Price
Peter	Kindle Fire HDX	129.99	Learning NoSQL	28.99
(timestamp)	(timestamp)	(timestamp)	(timestamp)	(timestamp)

Wide rows (2 billion cols)
Column names are data

Invoice
002

Buyer	X_7662_Name	X_7662_Price
Yoann	French cuisine	32.88
(timestamp)	(timestamp)	(timestamp)

*Keyspace = Database
Column Family ≈ Table*

```
set ColumnFamily['rowkey']['column'] = 'value';

get ColumnFamily['rowkey'];
get ColumnFamily['rowkey']['column'];

list ColumnFamily;

del ColumnFamily['rowkey'];
del ColumnFamily['rowkey']['column'];
```

KeySpace

```
CREATE KEYSPACE Excelsior WITH REPLICATION ={ 'class' :  
'SimpleStrategy', 'replication_factor' : 3 };
```

```
CREATE KEYSPACE "Excalibur" WITH REPLICATION ={ 'class' :  
'NetworkTopologyStrategy', 'dc1' : 3, 'dc2' : 2};
```

Tables

```
CREATE TABLE users  
  
( id int PRIMARY KEY, first_name text,  
last_name text, age int, hobbies list<test>);  
  
CREATE TABLE users  
  
( id int, first_name text, last_name text,  
age int, hobbies list<test>, PRIMARY KEY(id));
```

Cassandra CQL (KEYS)

No

- NO joins
- NO subqueries
- NO Group By functionality
- NO aggregation functions

Use

- Data denormalization
- Data duplication
- PK, Indexes
- Counters,
- CAS Transactions, Batches
- Advanced Data Structures

Keys Types

- Simple primary key

PRIMARY KEY (empID)

- Compound primary key

PRIMARY KEY (empID, subdeptID, depID)

- Composite primary key

PRIMARY KEY ((block_id, breed), color)

User-defined types

```
CREATE TYPE address (
    street text,
    city text,
    zip_code int,
    phones set<text> );
```

User-defined types

```
CREATE INDEX state_key ON users (state);
```

```
SELECT * FROM users WHERE gender = 'f' AND state = 'TX'  
ALLOW FILTERING;
```

Tables

```
CREATE TABLE users (
    user_id text PRIMARY KEY,
    first_name text, last_name text,
    emails set<text> );
INSERT INTO users (user_id, first_name, last_name, emails)
VALUES ('frodo', 'Frodo', 'Baggins',
        {'f@baggins.com', 'baggins@gmail.com'});
```

```
INSERT INTO clicks ( userid, url, date, name) VALUES (
3715e600-2eb0-11e2-81c1-0800200c9a66, 'http://apache.org',
'2013-10-09', 'Mary')

USING TTL 86400;
```

Ordering

```
CREATE TABLE timeseries (
    event_type text,
    insertion_time timestamp,
    PRIMARY KEY (event_type, insertion_time))
    WITH CLUSTERING ORDER BY (insertion_time DESC);

SELECT * FROM emp

WHERE empID IN (130,104) ORDER BY deptID DESC;
```

Slicing

//To retrieve events for the 12th of January 2014 between
3:50:00 and 4:37:30:

```
SELECT * FROM timeline WHERE day='12 Jan 2014' AND (hour,  
min) >= (3, 50) AND (hour, min, sec) <= (4, 37, 30);
```

Batching

```
BEGIN BATCH
```

```
INSERT INTO purchases (user, balance) VALUES ('user1', -8)
```

```
IF NOT EXISTS;
```

```
INSERT INTO purchases (user, expense_id, amount,  
description, paid) VALUES ('user1', 1, 8, 'bur', false);
```

```
APPLY BATCH;
```

Counters

```
CREATE TABLE counters.page_view_counts (counter_value  
counter, url_name varchar, page_name varchar, PRIMARY KEY  
(url_name, page_name) );
```

```
UPDATE counters.page_view_counts SET counter_value =  
counter_value + 1 WHERE url_name='www.datastax.com' AND  
page_name='home';
```

Lightweight Transactions

```
INSERT INTO users (login, email, name, login_count) VALUES  
('jdoe', 'jdoe@abc.com', 'Jane Doe', 1) IF NOT EXISTS;  
  
UPDATE users SET email = 'janedoe@abc.com' WHERE login =  
'jdoe' IF email = 'jdoe@abc.com';
```

DataStax Java Driver

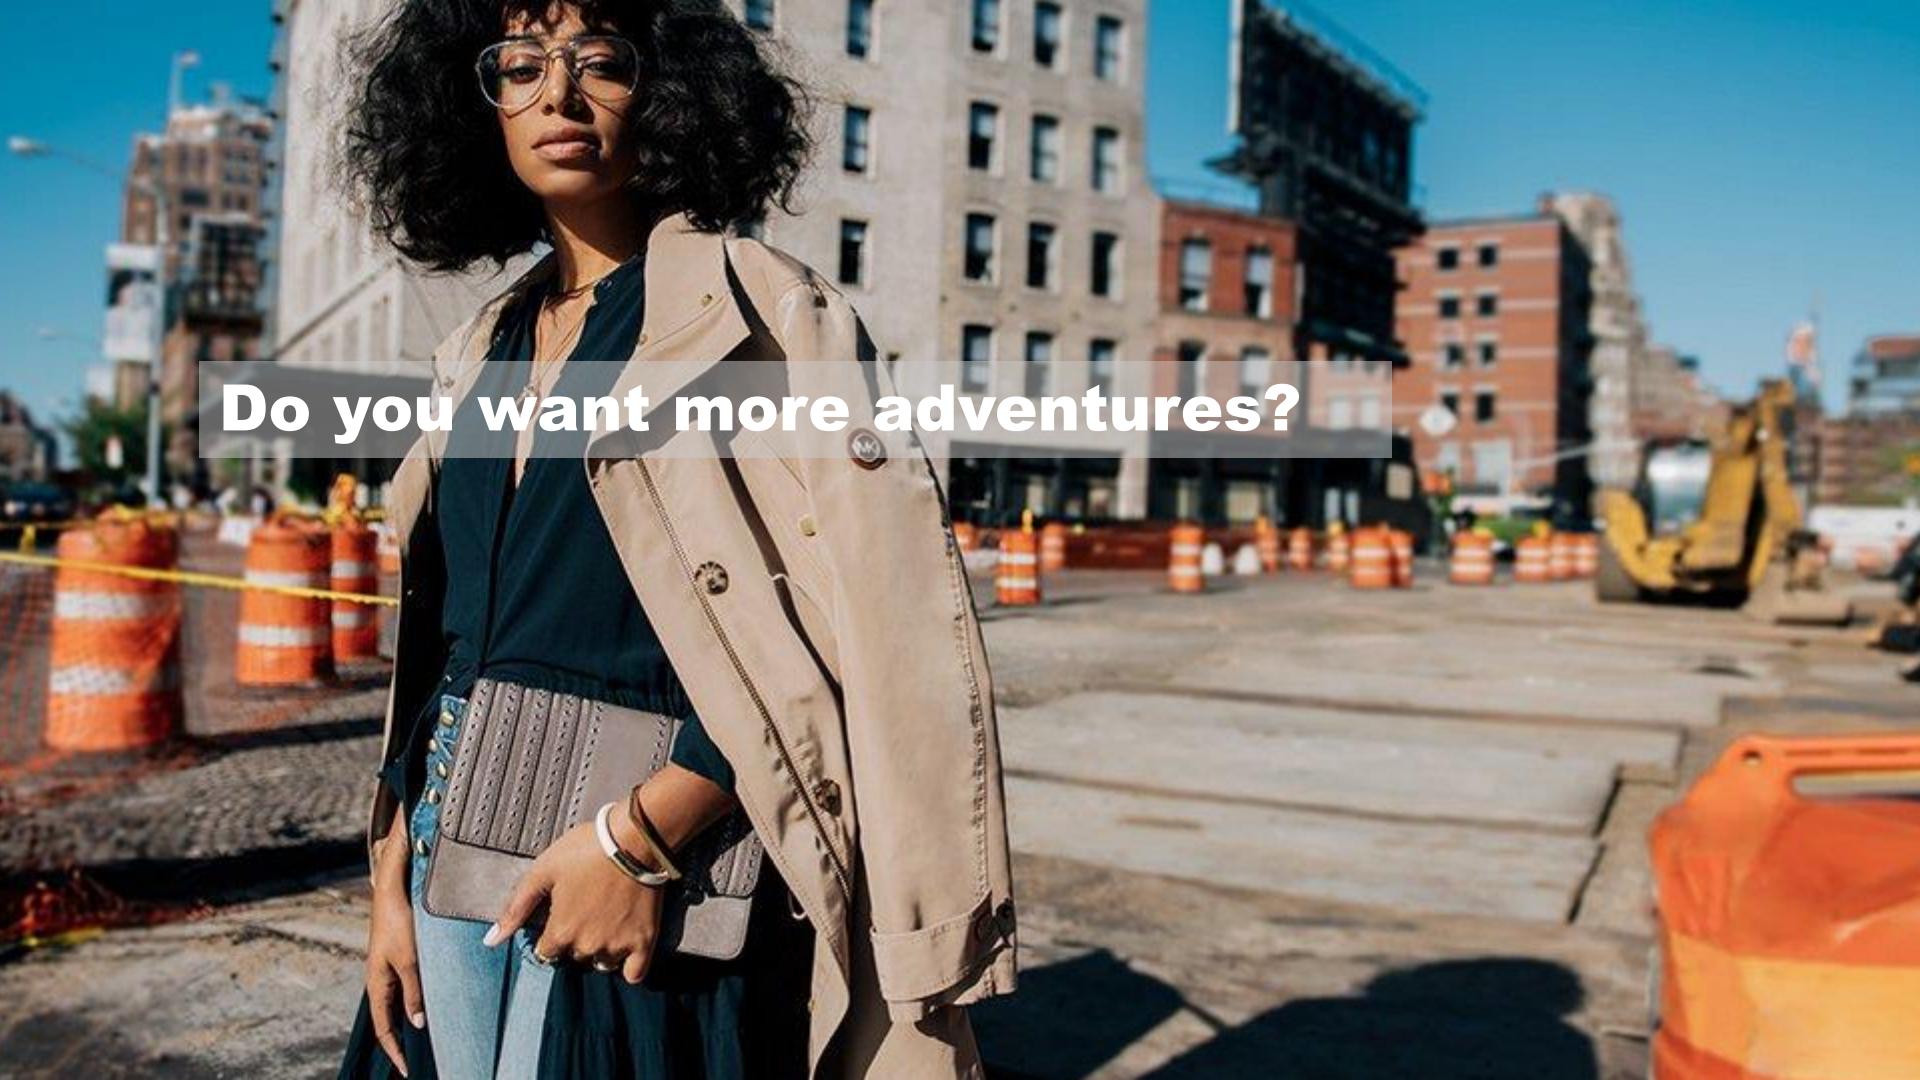
- DataStax developed a new protocol that doesn't have RPC limitations (Asynchronous I/O)
- Low-level API with simple mapping
- Works with CQL3
- QueryBuilder reminds CriteriaAPI
- Accessor-annotated interfaces

Async Query

```
public class AsynchronousExample extends SimpleClient {  
    public AsynchronousExample() {  
    }  
  
    public ResultSetFuture getRows() {  
        Query query = QueryBuilder.select().all().from("simplex", "songs");  
        return getSession().executeAsync(query);  
    }  
  
    public static void main(String[] args) {  
        AsynchronousExample client = new AsynchronousExample();  
        client.connect("127.0.0.1");  
        client.createSchema();  
        client.loadData();  
        ResultSetFuture results = client.getRows();  
        results.getUninterruptibly();  
        for (Row row : results)  
            System.out.printf("%s: %s / %s\n",  
                row.getString("artist"),  
                row.getString("title"),  
                row.getString("album"));  
    }  
    client.dropSchema("simplex");  
    client.close();  
}  
}
```

Accessor

```
@Accessor  
public interface UserAccessor {  
    @Query("SELECT * FROM complex.users WHERE id = :id")  
    User getUserNamed(@Param("userId") UUID id);  
  
    @Query("SELECT * FROM complex.users WHERE id = ?")  
    User getOnePosition(UUID userId);  
  
    @Query("UPDATE complex.users SET addresses[:name]=:  
    ResultSet addAddress(@Param("id") UUID id, @Param("'  
    @Query("SELECT * FROM complex.users")  
    public Result<User> getAll();  
  
    @Query("SELECT * FROM complex.users")  
    public ListenableFuture<Result<User>> getAllAsyn
```

A woman with dark curly hair and glasses, wearing a tan trench coat over a teal top and jeans, stands in an urban setting. She is positioned in front of a row of orange traffic cones and a yellow construction vehicle. In the background, there are several multi-story brick buildings under a clear blue sky.

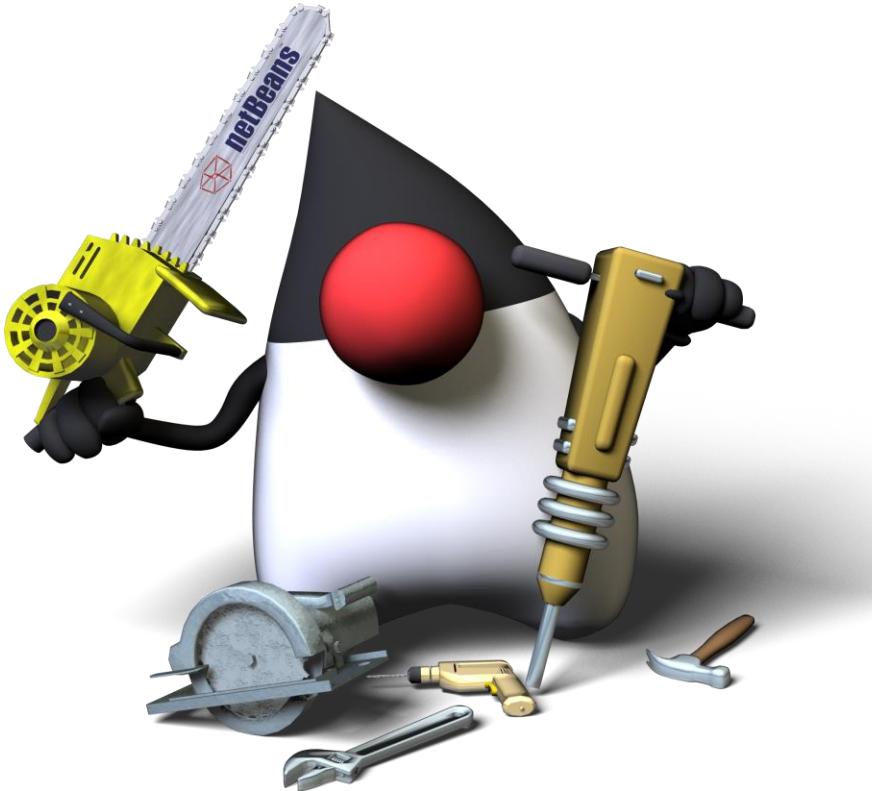
Do you want more adventures?

Other Cassandra's OM

- Achilles : well documented and provides transactions
- Astyanax : connection pool, thread safety and pagination
- Pelops : old project, good bycicle
- PlayORM : strange but powerful thing
- Easy-Cassandra : simple annotations + CRUD
- Thrift as low level API

MONGO WORLD

Case #7 : Customer wants MongoDB



Mongo Driver

```
DBCollection table = db.getCollection("user");
BasicDBObject document = new BasicDBObject();
document.put("name", "alex");
document.put("age", 27);
document.put("createdDate", new Date());
table.insert(document);
```

```
BasicDBObject query = new BasicDBObject();
query.put("name", "alex");

BasicDBObject newDocument = new BasicDBObject();
newDocument.put("name", "alex-updated");
newDocument.put("age", 28);

BasicDBObject updateObj = new BasicDBObject();
updateObj.put("$set", newDocument);
table.update(query, updateObj);
```

A woman with dark curly hair and glasses, wearing a tan trench coat over a teal top and jeans, stands in an urban setting. She is positioned in front of a row of orange traffic cones and a yellow construction vehicle. In the background, there are several multi-story brick buildings under a clear blue sky.

Do you have an ORM framework?

Morphia Entity

```
@Entity("employees")
class Employee {
    @Id ObjectId id;
    Address address;
    Key<Employee> manager;
    @Reference List<Employee> underlings = new ArrayList<
    @Property("left") Date endDate;
    //fields can be indexed for better performance
    @Indexed boolean active = false;
    //fields can loaded, but not saved
    @NotSaved String readButNotStored;
    //Lifecycle methods -- Pre/PostLoad, Pre/PostPersist.
    @PostLoad void postLoad(DBObject dbObj) { ... }
```

Morphia Entity

```
@EntityListeners(BackAccountWatcher.class)
public class BankAccount {
    @Id String id;
    Date lastUpdated = new Date();

    class BankAccountWatcher{
        @PrePersist void prePersist(BankAccount act) {act.lastUpdated
    }

    ... or

    class BankAccount {
        @Id String id;
        Date lastUpdated = new Date();
        @PrePersist void prePersist() {lastUpdated = new Date();}
    }
}
```

Other Mongo's OM

- Jongo : mongo - shell queries in Java-code
- EclipseLink : different support of different NoSQL databases
- MJORM : Google Code, XML mapping + MQL (SQL syntax for Mongo data extracting)
- DataNucleus : support many Js as JDO, JPA

Modern Web App



Polyglot Persistance

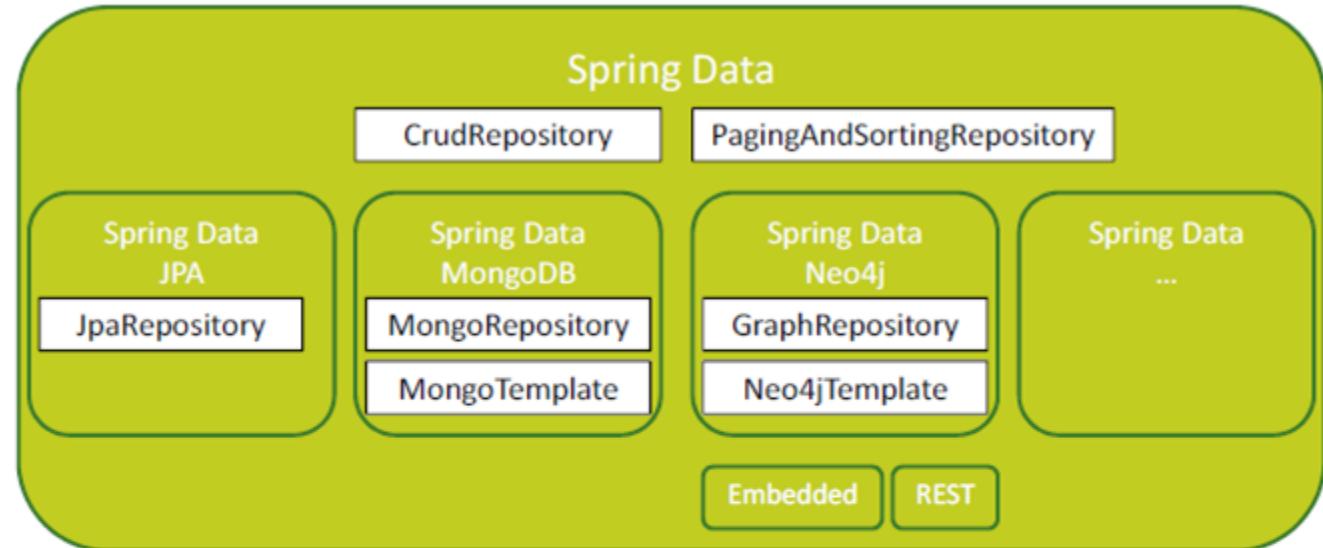
- Redis: Rapid access for reads and writes. No need to be durable
- RDBMS: Needs transactional updates and has tabular structure.
- Riak: Needs high availability across multiple locations.
Can merge inconsistent writes

Polyglot Persistance

- Neo4j: Rapidly traverse links between friends and ratings.
- MongoDB: Lots of reads, infrequent writes. Powerful aggregation mechanism.
- Cassandra: Large-scale analytics on large cluster. High volume of writes on multiple nodes

SPRING INTEGRATION

Spring Data



Spring Data MongoDB

- Templating : connection configs, collection lifecycle (create, drop), Map/Reduce + Aggregation
- Mapping: @Document, @Index, @Field
- Repository support: geospatial queries, queries derived from method signatures (at runtime)
- Paging, sorting, CRUD operations

Spring & MongoDB

```
public class Order {  
    @Id private String id;  
    private Date date;  
    @Field("custInfo") private String customerInfo;  
}  
  
public class Item {  
    private int quantity;  
    private double price;  
    @Field("desc") private String description;  
}  
  
public interface OrderRepository extends MongoRepository  
{  
    List<Order> findByItemsQuantity(int quantity);  
    List<Order> findByItemsPriceGreaterThanOrEqual(double price);  
}
```

Demo with JavaDriver 3.3



Best Practice: SQL for NoSQL





Any questions?